

Analyzing Police Activity with pandas

Chap 1: Preparing the data for analysis

In [1]:

Import plotting modules
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
plt.style.use('ggplot')

Stanford Open Policing Project dataset

In [2]:

Preparing the data

ri = pd.read_csv("datasets/police.csv")
ri.head()

Out[2]:

	state	stop_date	stop_time	county_name	driver_gender	driver_race	violation_raw	violation	search_conducted	search_type	stop_outcome	is_arrested	stop_duration	drugs_related_stop	district
0	RI	2005-01-04	12:55	NaN	M	White	Equipment/Inspection Violation	Equipment	False	NaN	Citation	False	0-15 Min	False	Zone X4
1	RI	2005-01-23	23:15	NaN	M	White	Speeding	Speeding	False	NaN	Citation	False	0-15 Min	False	Zone K3
2	RI	2005-02-17	04:15	NaN	M	White	Speeding	Speeding	False	NaN	Citation	False	0-15 Min	False	Zone X4
3	RI	2005-02-20	17:15	NaN	M	White	Call for Service	Other	False	NaN	Arrest Driver	True	16-30 Min	False	Zone X1
4	RI	2005-02-24	01:20	NaN	F	White	Speeding	Speeding	False	NaN	Citation	False	0-15 Min	False	Zone X3

In [3]:

ri.isnull().sum()

Out[3]:

state	0
stop_date	0
stop_time	0
county_name	91741
driver_gender	5205
driver_race	5202
violation_raw	5202
violation	5202
search_conducted	0
search_type	88434
stop_outcome	5202
is_arrested	5202
stop_duration	5202
drugs_related_stop	0
district	0

dtype: int64

In [4]:

ri.shape

Out[4]:

(91741, 15)

In [5]:

ri.drop('county_name', axis='columns', inplace=True)

In [6]:

ri.dropna(subset=['stop_date', 'stop_time']).shape

Out[6]:

(91741, 14)

No data is missing in the given columns.

Exercises

Examining the dataset

In [7]:

Import the pandas library as pd
import pandas as pd

Read 'police.csv' into a DataFrame named ri
ri = pd.read_csv('datasets/police.csv')

Examine the head of the DataFrame
print(ri.head(5))

Count the number of missing values in each column
print(ri.isnull().sum())

	state	stop_date	stop_time	county_name	driver_gender	driver_race	\
0	RI	2005-01-04	12:55	NaN	M	White	
1	RI	2005-01-23	23:15	NaN	M	White	
2	RI	2005-02-17	04:15	NaN	M	White	
3	RI	2005-02-20	17:15	NaN	M	White	
4	RI	2005-02-24	01:20	NaN	F	White	

	violation_raw	violation	search_conducted	search_type	\
0	Equipment/Inspection Violation	Equipment	False	NaN	
1	Speeding	Speeding	False	NaN	
2	Speeding	Speeding	False	NaN	
3	Call for Service	Other	False	NaN	
4	Speeding	Speeding	False	NaN	

	stop_outcome	is_arrested	stop_duration	drugs_related_stop	district
0	Citation	False	0-15 Min	False	Zone X4
1	Citation	False	0-15 Min	False	Zone K3
2	Citation	False	0-15 Min	False	Zone X4
3	Arrest Driver	True	16-30 Min	False	Zone X1
4	Citation	False	0-15 Min	False	Zone X3

state
0
stop_date
0
stop_time
0
county_name
91741
driver_gender
5205
driver_race
5202
violation_raw
5202
violation
5202
search_conducted
0
search_type
88434
stop_outcome
5202
is_arrested
5202
stop_duration
5202
drugs_related_stop
0
district
0
dtype: int64

Dropping columns

```
In [8]: # Count the number of missing values in each column
print(ri.isnull().sum())

# Examine the shape of the DataFrame
print(ri.shape)

# Drop the 'county_name' and 'state' columns
ri.drop(['county_name', 'state'], axis='columns', inplace=True)

# Examine the shape of the DataFrame (again)
print(ri.shape)
```

state	0
stop_date	0
stop_time	0
county_name	91741
driver_gender	5205
driver_race	5202
violation_raw	5202
violation	5202
search_conducted	0
search_type	88434
stop_outcome	5202
is_arrested	5202
stop_duration	5202
drugs_related_stop	0
district	0

dtype: int64
(91741, 15)
(91741, 13)

Dropping rows

```
In [9]: # Count the number of missing values in each column
print(ri.isnull().sum())

# Drop all rows that are missing 'driver_gender'
ri.dropna(subset=['driver_gender'], inplace=True)

# Count the number of missing values in each column (again)
print(ri.isnull().sum())

# Examine the shape of the DataFrame
print(ri.shape)
```

stop_date	0
stop_time	0
driver_gender	5205
driver_race	5202
violation_raw	5202
violation	5202
search_conducted	0
search_type	88434
stop_outcome	5202
is_arrested	5202
stop_duration	5202
drugs_related_stop	0
district	0

dtype: int64

stop_date	0
stop_time	0
driver_gender	0
driver_race	0
violation_raw	0
violation	0
search_conducted	0
search_type	83229
stop_outcome	0
is_arrested	0
stop_duration	0
drugs_related_stop	0
district	0

dtype: int64
(86536, 13)

Using proper data types

```
In [10]: # Examining the data types
ri.dtypes
```

```
Out[10]: stop_date      object
stop_time      object
driver_gender   object
driver_race     object
violation_raw   object
violation       object
search_conducted    bool
search_type      object
stop_outcome     object
is_arrested      object
stop_duration    object
drugs_related_stop    bool
district         object
dtype: object
```

```
In [11]: # Fixing a data type
apple = pd.read_csv("datasets/aapl_ohlcv.csv", usecols=[0,1,2], dtype='O')
apple.head(3)
```

	date	time	price
0	2/1/2018	16:00	170.1600037
1	3/1/2018	16:00	172.5299988
2	4/1/2018	16:00	172.5399933

```
In [12]: apple.price.dtype
```

```
Out[12]: dtype('O')
```

```
In [13]: apple['price'] = apple.price.astype('float')
apple.price.dtype
```

```
Out[13]: dtype('float64')
```

Exercises

Finding an incorrect data type

```
In [14]: ri.dtypes
```

```
Out[14]: stop_date      object
stop_time      object
driver_gender   object
driver_race     object
violation_raw   object
violation       object
search_conducted    bool
search_type      object
stop_outcome     object
is_arrested      object
stop_duration    object
drugs_related_stop    bool
district         object
dtype: object
```

```
In [15]: ri.head(3)
```

Out[15]:

	stop_date	stop_time	driver_gender	driver_race	violation_raw	violation	search_conducted	search_type	stop_outcome	is_arrested	stop_duration	drugs_related_stop	district
0	2005-01-04	12:55	M	White	Equipment/Inspection Violation	Equipment	False	NaN	Citation	False	0-15 Min	False	Zone X4
1	2005-01-23	23:15	M	White	Speeding	Speeding	False	NaN	Citation	False	0-15 Min	False	Zone K3
2	2005-02-17	04:15	M	White	Speeding	Speeding	False	NaN	Citation	False	0-15 Min	False	Zone X4

is_arrested should have a data type of bool

Fixing a data type

```
In [16]: # Examine the head of the 'is_arrested' column
print(ri.is_arrested.head())

# Check the data type of 'is_arrested'
print(ri.is_arrested.dtype)

# Change the data type of 'is_arrested' to 'bool'
ri['is_arrested'] = ri.is_arrested.astype('bool')

# Check the data type of 'is_arrested' (again)
print(ri.is_arrested.dtype)

0    False
1    False
2    False
3     True
4    False
Name: is_arrested, dtype: object
object
bool
```

Creating a DatetimeIndex

```
In [17]: # Using datetime format
ri.head(3)
```

Out[17]:

	stop_date	stop_time	driver_gender	driver_race	violation_raw	violation	search_conducted	search_type	stop_outcome	is_arrested	stop_duration	drugs_related_stop	district
0	2005-01-04	12:55	M	White	Equipment/Inspection Violation	Equipment	False	NaN	Citation	False	0-15 Min	False	Zone X4
1	2005-01-23	23:15	M	White	Speeding	Speeding	False	NaN	Citation	False	0-15 Min	False	Zone K3
2	2005-02-17	04:15	M	White	Speeding	Speeding	False	NaN	Citation	False	0-15 Min	False	Zone X4

```
In [18]: ri.dtypes
```

Out[18]:

stop_date object
stop_time object
driver_gender object
driver_race object
violation_raw object
violation object
search_conducted bool
search_type object
stop_outcome object
is_arrested bool
stop_duration object
drugs_related_stop bool
district object
dtype: object

```
In [19]: apple.head(3)
```

Out[19]:

	date	time	price
0	2/1/2018	16:00	170.160004
1	3/1/2018	16:00	172.529999
2	4/1/2018	16:00	172.539993

```
In [20]: apple.date.str.replace('/', '-').head(3)
```

Out[20]:

0 2-1-2018
1 3-1-2018
2 4-1-2018
Name: date, dtype: object

```
In [21]: combined = apple.date.str.cat(apple.time, sep=' ')
combined.head(3)
```

Out[21]:

0 2/1/2018 16:00
1 3/1/2018 16:00
2 4/1/2018 16:00
Name: date, dtype: object

```
In [22]: # Converting to datetime format
apple['date_and_time'] = pd.to_datetime(combined)
apple.head(3)
```

Out[22]:

	date	time	price	date_and_time
0	2/1/2018	16:00	170.160004	2018-02-01 16:00:00
1	3/1/2018	16:00	172.529999	2018-03-01 16:00:00
2	4/1/2018	16:00	172.539993	2018-04-01 16:00:00

```
In [23]: apple.dtypes
```

Out[23]:

date object
time object
price float64
date_and_time datetime64[ns]
dtype: object

```
In [24]: apple.set_index('date_and_time', inplace=True)
apple.head(3)
```

Out[24]:

	date	time	price
date_and_time			
2018-02-01 16:00:00	2/1/2018	16:00	170.160004
2018-03-01 16:00:00	3/1/2018	16:00	172.529999
2018-04-01 16:00:00	4/1/2018	16:00	172.539993

```
In [25]: ▶ apple.index

Out[25]: DatetimeIndex(['2018-02-01 16:00:00', '2018-03-01 16:00:00',
                        '2018-04-01 16:00:00', '2018-05-01 16:00:00',
                        '2018-08-01 16:00:00', '2018-09-01 16:00:00',
                        '2018-10-01 16:00:00', '2018-11-01 16:00:00',
                        '2018-12-01 16:00:00', '2018-01-16 16:00:00',
                        '2018-01-17 16:00:00', '2018-01-18 16:00:00',
                        '2018-01-19 16:00:00', '2018-01-22 16:00:00',
                        '2018-01-23 16:00:00', '2018-01-24 16:00:00',
                        '2018-01-25 16:00:00', '2018-01-26 16:00:00',
                        '2018-01-29 16:00:00', '2018-01-30 16:00:00',
                        '2018-01-31 16:00:00', '2018-01-02 16:00:00',
                        '2018-02-02 16:00:00', '2018-05-02 16:00:00',
                        '2018-06-02 16:00:00', '2018-07-02 16:00:00',
                        '2018-08-02 16:00:00', '2018-09-02 16:00:00',
                        '2018-12-02 16:00:00', '2018-02-13 16:00:00',
                        '2018-02-14 16:00:00', '2018-02-15 16:00:00',
                        '2018-02-16 16:00:00', '2018-02-20 16:00:00',
                        '2018-02-21 16:00:00', '2018-02-22 16:00:00',
                        '2018-02-23 16:00:00', '2018-02-26 16:00:00',
                        '2018-02-27 16:00:00', '2018-02-28 16:00:00',
                        '2018-01-03 16:00:00', '2018-02-03 16:00:00',
                        '2018-05-03 16:00:00', '2018-06-03 16:00:00',
                        '2018-07-03 16:00:00', '2018-08-03 16:00:00',
                        '2018-09-03 16:00:00', '2018-12-03 16:00:00',
                        '2018-03-13 16:00:00', '2018-03-14 16:00:00',
                        '2018-03-15 16:00:00', '2018-03-16 16:00:00',
                        '2018-03-19 16:00:00', '2018-03-20 16:00:00',
                        '2018-03-21 16:00:00', '2018-03-22 16:00:00',
                        '2018-03-23 16:00:00', '2018-03-26 16:00:00',
                        '2018-03-27 16:00:00', '2018-03-28 16:00:00',
                        '2018-03-29 16:00:00'],
                        dtype='datetime64[ns]', name='date_and_time', freq=None)

In [26]: ▶ apple.columns

Out[26]: Index(['date', 'time', 'price'], dtype='object')
```

Exercises

Combining object columns

```
In [27]: ▶ # Concatenate 'stop_date' and 'stop_time' (separated by a space)
combined = ri.stop_date.str.cat(ri.stop_time, sep=' ')

# Convert 'combined' to datetime format
ri['stop_datetime'] = pd.to_datetime(combined)

# Examine the data types of the DataFrame
print(ri.dtypes)

stop_date          object
stop_time          object
driver_gender      object
driver_race        object
violation_raw      object
violation          object
search_conducted   bool
search_type        object
stop_outcome       object
is_arrested        bool
stop_duration      object
drugs_related_stop bool
district           object
stop_datetime      datetime64[ns]
dtype: object
```

Setting the index

```
In [28]: ▶ # Set 'stop_datetime' as the index
ri.set_index('stop_datetime', inplace=True)

# Examine the index
print(ri.index)

# Examine the columns
print(ri.columns)

DatetimeIndex(['2005-01-04 12:55:00', '2005-01-23 23:15:00',
                '2005-02-17 04:15:00', '2005-02-20 17:15:00',
                '2005-02-24 01:20:00', '2005-03-14 10:00:00',
                '2005-03-29 21:55:00', '2005-04-04 21:25:00',
                '2005-07-14 11:20:00', '2005-07-14 19:55:00',
                ...,
                '2015-12-31 13:23:00', '2015-12-31 18:59:00',
                '2015-12-31 19:13:00', '2015-12-31 20:20:00',
                '2015-12-31 20:50:00', '2015-12-31 21:21:00',
                '2015-12-31 21:59:00', '2015-12-31 22:04:00',
                '2015-12-31 22:09:00', '2015-12-31 22:47:00'],
                dtype='datetime64[ns]', name='stop_datetime', length=86536, freq=None)
Index(['stop_date', 'stop_time', 'driver_gender', 'driver_race',
       'violation_raw', 'violation', 'search_conducted', 'search_type',
       'stop_outcome', 'is_arrested', 'stop_duration', 'drugs_related_stop',
       'district'],
      dtype='object')
```

Chap 2: Exploring the relationship between gender and policing

```
In [29]: ▶ # Import plotting modules
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
plt.style.use('ggplot')
```

Do the genders commit different violations?

```
In [30]: ▶ # Counting unique values
ri.stop_outcome.value_counts()

Out[30]: Citation          77091
Warning                   5136
Arrest Driver             2735
No Action                  624
N/D                        607
Arrest Passenger          343
Name: stop_outcome, dtype: int64

In [31]: ▶ ri.stop_outcome.value_counts().sum()

Out[31]: 86536

In [32]: ▶ ri.shape

Out[32]: (86536, 13)

In [33]: ▶ # Expressing counts as proportions
77091/86536

Out[33]: 0.8908546731995932
```

```
In [34]: ri.stop_outcome.value_counts(normalize=True)

Out[34]: Citation      0.890855
Warning      0.059351
Arrest Driver  0.031605
No Action     0.007211
N/D           0.007014
Arrest Passenger 0.003964
Name: stop_outcome, dtype: float64

In [35]: # Filtering DataFrame rows
ri.driver_race.value_counts()

Out[35]: White      61870
Black      12285
Hispanic    9727
Asian      2389
Other       265
Name: driver_race, dtype: int64

In [36]: white = ri[ri.driver_race == 'White']
white.shape

Out[36]: (61870, 13)

In [37]: # Comparing stop outcomes for two groups
white.stop_outcome.value_counts(normalize=True)

Out[37]: Citation      0.902263
Warning      0.057508
Arrest Driver  0.024018
No Action     0.007031
N/D           0.006433
Arrest Passenger 0.002748
Name: stop_outcome, dtype: float64

In [38]: asian = ri[ri.driver_race == 'Asian']
asian.stop_outcome.value_counts(normalize=True)

Out[38]: Citation      0.922980
Warning      0.045207
Arrest Driver  0.017581
No Action     0.008372
N/D           0.004186
Arrest Passenger 0.001674
Name: stop_outcome, dtype: float64
```

Exercises

Examining traffic violations

```
In [39]: # Count the unique values in 'violation'
print(ri.violation.value_counts())

# Express the counts as proportions
print(ri.violation.value_counts(normalize=True))

Speeding      48423
Moving violation 16224
Equipment     10921
Other         4409
Registration/plates 3703
Seat belt     2856
Name: violation, dtype: int64
Speeding      0.559571
Moving violation 0.187483
Equipment     0.126202
Other         0.050950
Registration/plates 0.042791
Seat belt     0.033004
Name: violation, dtype: float64
```

Comparing violations by gender

```
In [40]: # Create a DataFrame of female drivers
female = ri[ri.driver_gender == 'F']

# Create a DataFrame of male drivers
male = ri[ri.driver_gender == 'M']

# Compute the violations by female drivers (as proportions)
print(female.violation.value_counts(normalize=True))

# Compute the violations by male drivers (as proportions)
print(male.violation.value_counts(normalize=True))

Speeding      0.658114
Moving violation 0.138218
Equipment     0.105199
Registration/plates 0.044418
Other         0.029738
Seat belt     0.024312
Name: violation, dtype: float64
Speeding      0.522243
Moving violation 0.206144
Equipment     0.134158
Other         0.058985
Registration/plates 0.042175
Seat belt     0.036296
Name: violation, dtype: float64
```

Does gender affect who gets a ticket for speeding?

```
In [41]: # Filtering by multiple conditions
female = ri[ri.driver_gender == 'F']
female.shape

Out[41]: (23774, 13)

In [42]: female_and_arrested = ri[(ri.driver_gender == 'F') &
                                   (ri.is_arrested == True)]
female_and_arrested.shape

Out[42]: (669, 13)

In [43]: female_or_arrested = ri[(ri.driver_gender == 'F') |
                                   (ri.is_arrested == True)]
female_or_arrested.shape

Out[43]: (26183, 13)
```

Exercise

Filtering by multiple conditions

Which one of these commands would filter the ri DataFrame to only include female drivers who were stopped for a speeding violation?

```
In [44]: ri[(ri.driver_gender == 'F') & (ri.violation == 'Speeding')].head(3)
```

Out[44]:

	stop_date	stop_time	driver_gender	driver_race	violation_raw	violation	search_conducted	search_type	stop_outcome	is_arrested	stop_duration	drugs_related_stop	district
	stop_datetime												
	2005-02-24 01:20:00	2005-02-24 01:20	F	White	Speeding	Speeding	False	NaN	Citation	False	0-15 Min	False	Zone X3
	2005-03-14 10:00:00	2005-03-14 10:00	F	White	Speeding	Speeding	False	NaN	Citation	False	0-15 Min	False	Zone K3
	2005-07-14 11:20:00	2005-07-14 11:20	F	White	Speeding	Speeding	False	NaN	Citation	False	0-15 Min	False	Zone X4

Comparing speeding outcomes by gender

```
In [45]: # Create a DataFrame of female drivers stopped for speeding
female_and_speeding = ri[(ri.driver_gender == 'F') & (ri.violation == 'Speeding')]

# Create a DataFrame of male drivers stopped for speeding
male_and_speeding = ri[(ri.driver_gender == 'M') & (ri.violation == 'Speeding')]

# Compute the stop outcomes for female drivers (as proportions)
print(female_and_speeding.stop_outcome.value_counts(normalize=True))

# Compute the stop outcomes for male drivers (as proportions)
print(male_and_speeding.stop_outcome.value_counts(normalize=True))
```

```
Citation      0.952192
Warning       0.040074
Arrest Driver 0.005752
N/D           0.000959
Arrest Passenger 0.000639
No Action     0.000383
Name: stop_outcome, dtype: float64
Citation      0.944595
Warning       0.036184
Arrest Driver 0.015895
Arrest Passenger 0.001281
No Action     0.001068
N/D           0.000976
Name: stop_outcome, dtype: float64
```

Does gender affect whose vehicle is searched?

```
In [46]: ri.isnull().sum()
```

Out[46]:

stop_date	0
stop_time	0
driver_gender	0
driver_race	0
violation_raw	0
violation	0
search_conducted	0
search_type	83229
stop_outcome	0
is_arrested	0
stop_duration	0
drugs_related_stop	0
district	0

dtype: int64

```
In [47]: np.mean([0,1,0,0])
```

Out[47]: 0.25

```
In [48]: np.mean([False,True,False,False])
```

Out[48]: 0.25

```
In [49]: # Taking the mean of a Boolean Series
ri.is_arrested.value_counts(normalize=True)
```

Out[49]:

False	0.964431
True	0.035569

Name: is_arrested, dtype: float64

```
In [50]: ri.is_arrested.mean()
```

Out[50]: 0.0355690117407784

```
In [51]: ri.is_arrested.dtype
```

Out[51]: dtype('bool')

```
In [52]: # Comparing groups using groupby
ri.district.unique()
```

Out[52]: array(['Zone X4', 'Zone K3', 'Zone X1', 'Zone X3', 'Zone K1', 'Zone K2'], dtype=object)

```
In [53]: ri[ri.district == 'Zone K1'].is_arrested.mean()
```

Out[53]: 0.024349083895853423

```
In [54]: ri[ri.district == 'Zone K2'].is_arrested.mean()
```

Out[54]: 0.030800588834786546

```
In [55]: ri.groupby('district').is_arrested.mean()
```

Out[55]:

district	
Zone K1	0.024349
Zone K2	0.030801
Zone K3	0.032311
Zone X1	0.023494
Zone X3	0.034871
Zone X4	0.048038

Name: is_arrested, dtype: float64

```
In [56]: # Grouping by multiple categories
ri.groupby(['district', 'driver_gender']).is_arrested.mean()
```

Out[56]:

district	driver_gender	
Zone K1	F	0.019169
	M	0.026588
Zone K2	F	0.022196
	M	0.034285
Zone K3	F	0.025156
	M	0.034961
Zone X1	F	0.019646
	M	0.024563
Zone X3	F	0.027188
	M	0.038166
Zone X4	F	0.042149
	M	0.049956

Name: is_arrested, dtype: float64

In [57]: ri.groupby(['driver_gender','district']).is_arrested.mean()

Out[57]: driver_gender district
F Zone K1 0.019169
Zone K2 0.022196
Zone K3 0.025156
Zone X1 0.019646
Zone X3 0.027188
Zone X4 0.042149
M Zone K1 0.026588
Zone K2 0.034285
Zone K3 0.034961
Zone X1 0.024563
Zone X3 0.038166
Zone X4 0.049956
Name: is_arrested, dtype: float64

Exercises

Calculate the search rate

In [58]: # Check the data type of 'search_conducted'
print(ri.search_conducted.dtype)

Calculate the search rate by counting the values
print(ri.search_conducted.value_counts(normalize=True))

Calculate the search rate by taking the mean
print(ri.search_conducted.mean())

bool
False 0.961785
True 0.038215
Name: search_conducted, dtype: float64
0.0382153092354627

Comparing search rates by gender

In [59]: # Calculate the search rate for female drivers
print(ri[ri.driver_gender == 'F'].search_conducted.mean())

0.019180617481282074

In [60]: # Calculate the search rate for male drivers
print(ri[ri.driver_gender == 'M'].search_conducted.mean())

0.04542557598546892

In [61]: # Calculate the search rate for both groups simultaneously
print(ri.groupby('driver_gender').search_conducted.mean())

driver_gender
F 0.019181
M 0.045426
Name: search_conducted, dtype: float64

Adding a second factor to the analysis

In [62]: # Calculate the search rate for each combination of gender and violation
print(ri.groupby(['driver_gender','violation']).search_conducted.mean())

driver_gender violation
F Equipment 0.039984
Moving violation 0.039257
Other 0.041018
Registration/plates 0.054924
Seat belt 0.017301
Speeding 0.008309
M Equipment 0.071496
Moving violation 0.061524
Other 0.046191
Registration/plates 0.108802
Seat belt 0.035119
Speeding 0.027885
Name: search_conducted, dtype: float64

In [63]: # Reverse the ordering to group by violation before gender
print(ri.groupby(['violation','driver_gender']).search_conducted.mean())

violation driver_gender
Equipment F 0.039984
M 0.071496
Moving violation F 0.039257
M 0.061524
Other F 0.041018
M 0.046191
Registration/plates F 0.054924
M 0.108802
Seat belt F 0.017301
M 0.035119
Speeding F 0.008309
M 0.027885
Name: search_conducted, dtype: float64

Does gender affect who is frisked during a search?

In [64]: # Examining the search types
ri.search_conducted.value_counts()

Out[64]: False 83229
True 3307
Name: search_conducted, dtype: int64

In [65]: ri.search_type.value_counts(dropna=False)

Out[65]: NaN 83229
Incident to Arrest 1290
Probable Cause 924
Inventory 219
Reasonable Suspicion 214
Protective Frisk 164
Incident to Arrest,Inventory 123
Incident to Arrest,Probable Cause 100
Probable Cause,Reasonable Suspicion 54
Probable Cause,Protective Frisk 35
Incident to Arrest,Inventory,Probable Cause 35
Incident to Arrest,Protective Frisk 33
Inventory,Probable Cause 25
Protective Frisk,Reasonable Suspicion 19
Incident to Arrest,Inventory,Protective Frisk 18
Incident to Arrest,Probable Cause,Protective Frisk 13
Inventory,Protective Frisk 12
Incident to Arrest,Reasonable Suspicion 8
Incident to Arrest,Probable Cause,Reasonable Suspicion 5
Probable Cause,Protective Frisk,Reasonable Suspicion 5
Incident to Arrest,Inventory,Reasonable Suspicion 4
Incident to Arrest,Protective Frisk,Reasonable Suspicion 2
Inventory,Reasonable Suspicion 2
Inventory,Protective Frisk,Reasonable Suspicion 1
Inventory,Probable Cause,Reasonable Suspicion 1
Inventory,Probable Cause,Protective Frisk 1
Name: search_type, dtype: int64

value_counts() excludes missing values by default

dropna=False displays missing values

In [66]:

Searching for a string
ri['inventory'] = ri.search_type.str.contains('Inventory', na=False)

In [67]:

ri.inventory.dtype

Out[67]:

dtype('bool')

In [68]:

ri.inventory.sum()

Out[68]:

441

In [69]:

Calculating the inventory rate
ri.inventory.mean()

Out[69]:

0.0050961449570121106

In [70]:

searched = ri[ri.search_conducted == True]
searched.inventory.mean()

Out[70]:

0.13335349259147264

0.5% of all traffic stops resulted in an inventory
13.3% of searches included an inventory

Exercises

Counting protective frisks

In [71]:

Count the 'search_type' values
print(ri.search_type.value_counts())

Incident to Arrest	1290
Probable Cause	924
Inventory	219
Reasonable Suspicion	214
Protective Frisk	164
Incident to Arrest,Inventory	123
Incident to Arrest,Probable Cause	100
Probable Cause,Reasonable Suspicion	54
Probable Cause,Protective Frisk	35
Incident to Arrest,Inventory,Probable Cause	35
Incident to Arrest,Protective Frisk	33
Inventory,Probable Cause	25
Protective Frisk,Reasonable Suspicion	19
Incident to Arrest,Inventory,Protective Frisk	18
Incident to Arrest,Probable Cause,Protective Frisk	13
Inventory,Protective Frisk	12
Incident to Arrest,Reasonable Suspicion	8
Probable Cause,Protective Frisk,Reasonable Suspicion	5
Incident to Arrest,Probable Cause,Reasonable Suspicion	5
Incident to Arrest,Inventory,Reasonable Suspicion	4
Inventory,Reasonable Suspicion	2
Incident to Arrest,Protective Frisk,Reasonable Suspicion	2
Inventory,Protective Frisk,Reasonable Suspicion	1
Inventory,Probable Cause,Reasonable Suspicion	1
Inventory,Probable Cause,Protective Frisk	1

Name: search_type, dtype: int64

In [72]:

Check if 'search_type' contains the string 'Protective Frisk'
ri['frisk'] = ri.search_type.str.contains('Protective Frisk', na=False)

Check the data type of 'frisk'
print(ri.frisk.dtype)

bool

In [73]:

Take the sum of 'frisk'
print(ri.frisk.sum())

303

Comparing frisk rates by gender

In [74]:

Create a DataFrame of stops in which a search was conducted
searched = ri[ri.search_conducted == True]

Calculate the overall frisk rate by taking the mean of 'frisk'
print(searched.frisk.mean())

0.09162382824312065

In [75]:

Calculate the frisk rate for each gender
print(searched.groupby('driver_gender').frisk.mean())

driver_gender	
F	0.074561
M	0.094353

Name: frisk, dtype: float64

Chap 3: Visual exploratory data analysis

In [76]:

Import plotting modules
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
plt.style.use('ggplot')

Does time of day affect arrest rate?

In [92]:

apple = pd.read_csv('datasets/aapl_ohlcv.csv')
apple['date_and_time'] = pd.to_datetime(apple.date.str.replace('/', '-').str.cat(apple.time, sep=' '))
apple.drop(columns=['date', 'time'], inplace=True)

In [93]:

Accessing datetime attributes
apple.head(5)

Out[93]:

	price	volume	date_and_time
0	170.160004	25555900	2018-02-01 16:00:00
1	172.529999	29517900	2018-03-01 16:00:00
2	172.539993	22434600	2018-04-01 16:00:00
3	173.440002	23660000	2018-05-01 16:00:00
4	174.350006	20567800	2018-08-01 16:00:00

In [94]:

apple.dtypes

Out[94]:

price	float64
volume	int64
date_and_time	datetime64[ns]

dtype: object


```
In [95]: apple.date_and_time.dt.month.head(5)
```

```
Out[95]: 0    2
         1    3
         2    4
         3    5
         4    8
         Name: date_and_time, dtype: int64
```

```
In [96]: apple.set_index('date_and_time', inplace=True)
         apple.index
```

```
Out[96]: DatetimeIndex(['2018-02-01 16:00:00', '2018-03-01 16:00:00',
                        '2018-04-01 16:00:00', '2018-05-01 16:00:00',
                        '2018-08-01 16:00:00', '2018-09-01 16:00:00',
                        '2018-10-01 16:00:00', '2018-11-01 16:00:00',
                        '2018-12-01 16:00:00', '2018-01-16 16:00:00',
                        '2018-01-17 16:00:00', '2018-01-18 16:00:00',
                        '2018-01-19 16:00:00', '2018-01-22 16:00:00',
                        '2018-01-23 16:00:00', '2018-01-24 16:00:00',
                        '2018-01-25 16:00:00', '2018-01-26 16:00:00',
                        '2018-01-29 16:00:00', '2018-01-30 16:00:00',
                        '2018-01-31 16:00:00', '2018-01-02 16:00:00',
                        '2018-02-02 16:00:00', '2018-05-02 16:00:00',
                        '2018-06-02 16:00:00', '2018-07-02 16:00:00',
                        '2018-08-02 16:00:00', '2018-09-02 16:00:00',
                        '2018-12-02 16:00:00', '2018-02-13 16:00:00',
                        '2018-02-14 16:00:00', '2018-02-15 16:00:00',
                        '2018-02-16 16:00:00', '2018-02-20 16:00:00',
                        '2018-02-21 16:00:00', '2018-02-22 16:00:00',
                        '2018-02-23 16:00:00', '2018-02-26 16:00:00',
                        '2018-02-27 16:00:00', '2018-02-28 16:00:00',
                        '2018-01-03 16:00:00', '2018-02-03 16:00:00',
                        '2018-05-03 16:00:00', '2018-06-03 16:00:00',
                        '2018-07-03 16:00:00', '2018-08-03 16:00:00',
                        '2018-09-03 16:00:00', '2018-12-03 16:00:00',
                        '2018-03-13 16:00:00', '2018-03-14 16:00:00',
                        '2018-03-15 16:00:00', '2018-03-16 16:00:00',
                        '2018-03-19 16:00:00', '2018-03-20 16:00:00',
                        '2018-03-21 16:00:00', '2018-03-22 16:00:00',
                        '2018-03-23 16:00:00', '2018-03-26 16:00:00',
                        '2018-03-27 16:00:00', '2018-03-28 16:00:00',
                        '2018-03-29 16:00:00'],
                        dtype='datetime64[ns]', name='date_and_time', freq=None)
```

```
In [97]: apple.index.month
```

```
Out[97]: Int64Index([ 2,  3,  4,  5,  8,  9, 10, 11, 12,  1,  1,  1,  1,  1,  1,  1,  1,
                    1,  1,  1,  1,  1,  2,  5,  6,  7,  8,  9, 12,  2,  2,  2,  2,  2,
                    2,  2,  2,  2,  2,  1,  2,  5,  6,  7,  8,  9, 12,  3,  3,  3,
                    3,  3,  3,  3,  3,  3,  3,  3],
                    dtype='int64', name='date_and_time')
```

```
In [98]: # Calculating the monthly mean price
         apple.price.mean()
```

Out[98]: 172.2736063508197


apple.groupby('month').price.mean() is invalid

```
In [99]: apple.groupby(apple.index.month).price.mean()
```

```
Out[99]: date_and_time
         1    174.189998
         2    171.511429
         3    173.956429
         4    172.539993
         5    169.250005
         6    166.370003
         7    169.014999
         8    170.039998
         9    169.860006
        10    173.160004
        11    174.589996
        12    171.656662
         Name: price, dtype: float64
```

```
In [100]: monthly_price = apple.groupby(apple.index.month).price.mean()
```

```
In [101]: # Plotting the monthly mean price
         monthly_price.plot();
         plt.xlabel('Month')
         plt.ylabel('Price')
         plt.title('Monthly mean stock price for Apple')
         plt.show()
```



Exercises

Calculating the hourly arrest rate

```
In [103]: # Calculate the overall arrest rate
         print(ri.is_arrested.mean())
```

0.0355690117407784

```
In [106]: # Calculate the hourly arrest rate
print(ri.is_arrested.groupby(ri.index.hour).mean())

# Save the hourly arrest rate
hourly_arrest_rate = ri.is_arrested.groupby(ri.index.hour).mean()

stop_datetime
0      0.051431
1      0.064932
2      0.060798
3      0.060549
4      0.048000
5      0.042781
6      0.013813
7      0.013032
8      0.021854
9      0.025206
10     0.028213
11     0.028897
12     0.037399
13     0.030776
14     0.030605
15     0.030679
16     0.035281
17     0.040619
18     0.038204
19     0.032245
20     0.038107
21     0.064541
22     0.048666
23     0.047592
Name: is_arrested, dtype: float64
```


Plotting the hourly arrest rate

```
In [110]: # Import matplotlib.pyplot as plt
import matplotlib.pyplot as plt

# Create a line plot of 'hourly_arrest_rate'
hourly_arrest_rate.plot()

# Add the xlabel, ylabel, and title
plt.xlabel('Hour')
plt.ylabel('Arrest Rate')
plt.title('Arrest Rate by Time of Day')

# Display the plot
plt.show()
```



Are drug-related stops on the rise?

```
In [111]: # Resampling the price
apple.head(5)
```

Out[111]:

	price	volume
date_and_time		
2018-02-01 16:00:00	170.160004	25555900
2018-03-01 16:00:00	172.529999	29517900
2018-04-01 16:00:00	172.539993	22434600
2018-05-01 16:00:00	173.440002	23660000
2018-08-01 16:00:00	174.350006	20567800

```
In [112]: apple.groupby(apple.index.month).price.mean().head(3)
```

Out[112]:

date_and_time	price
1	174.189998
2	171.511429
3	173.956429

Name: price, dtype: float64

```
In [113]: apple.price.resample('M').mean().head(3)
```

Out[113]:

date_and_time	price
2018-01-31	174.189998
2018-02-28	171.511429
2018-03-31	173.956429

Freq: M, Name: price, dtype: float64

```
In [115]: # Resampling the volume
apple.head(5)
```

Out[115]:

	price	volume
date_and_time		
2018-02-01 16:00:00	170.160004	25555900
2018-03-01 16:00:00	172.529999	29517900
2018-04-01 16:00:00	172.539993	22434600
2018-05-01 16:00:00	173.440002	23660000
2018-08-01 16:00:00	174.350006	20567800

```
In [116]: apple.volume.resample('M').mean().head(3)
```

Out[116]:

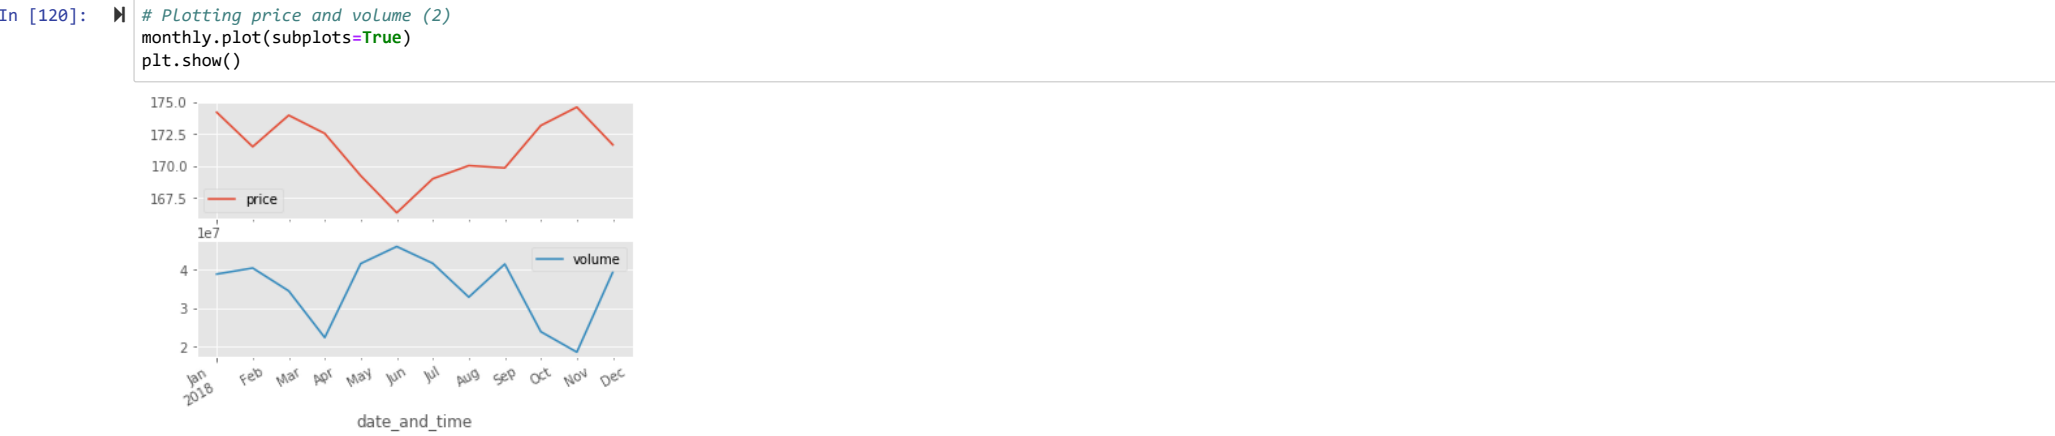

date_and_time	volume
2018-01-31	3.888188e+07
2018-02-28	4.044285e+07
2018-03-31	3.451521e+07

Freq: M, Name: volume, dtype: float64

```
In [118]: # Concatenating price and volume
monthly_price = apple.price.resample('M').mean()
monthly_volume = apple.volume.resample('M').mean()
monthly = pd.concat([monthly_price, monthly_volume], axis='columns')
monthly.head(3)
```

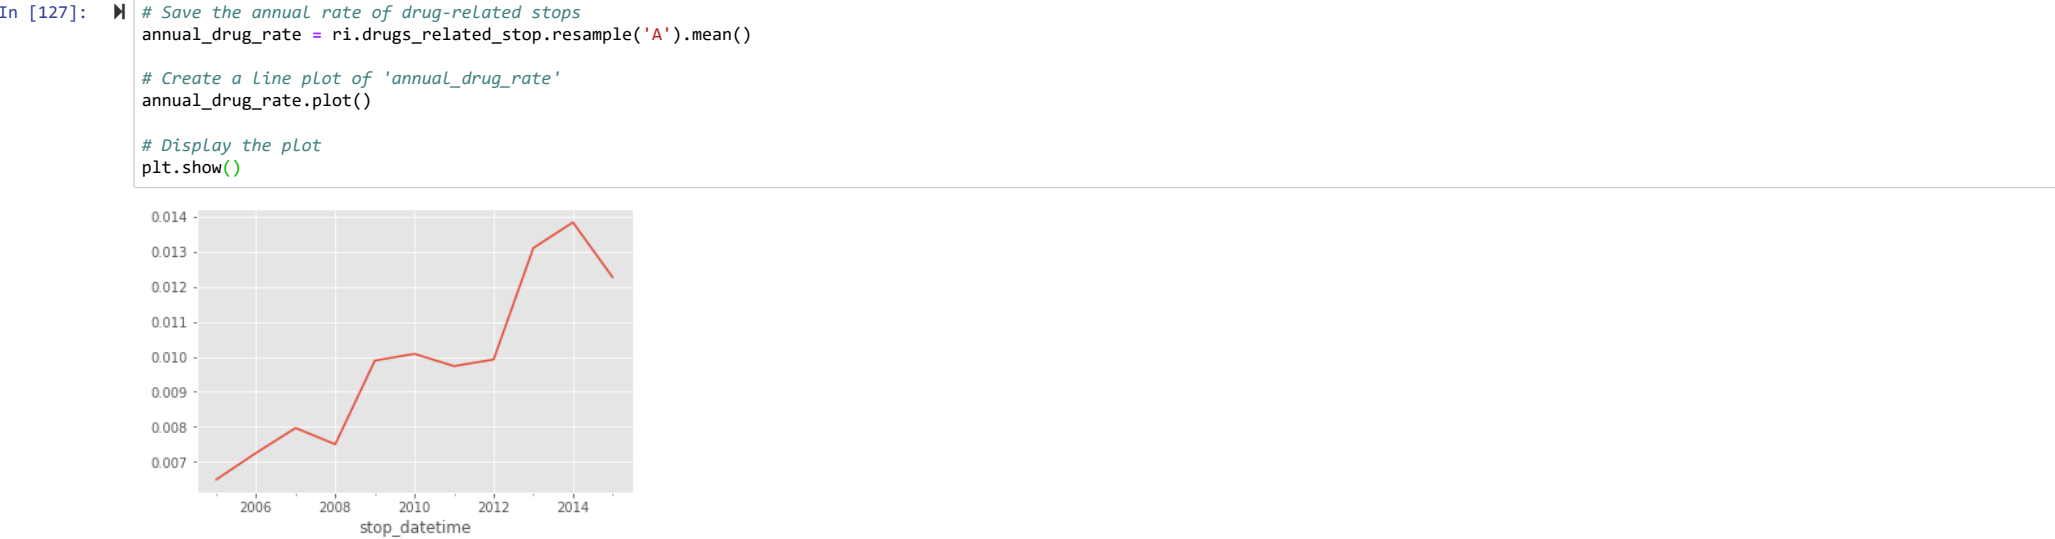

Out[118]:

	price	volume
date_and_time		
2018-01-31	174.189998	3.888188e+07
2018-02-28	171.511429	4.044285e+07
2018-03-31	173.956429	3.451521e+07



Exercises

Plotting drug-related stops



Comparing drug and search rates



What violations are caught in each district?

```
In [141]: # Computing a frequency table
table = pd.crosstab(ri.driver_race, ri.driver_gender)
table
```

Out[141]:

driver_gender	F	M
driver_race		
Asian	551	1838
Black	2681	9604
Hispanic	1953	7774
Other	53	212
White	18536	43334

Frequency table: Tally of how many times each combination of values occurs

```
In [142]: ri[(ri.driver_race == 'Asian') & (ri.driver_gender == 'F')].shape
```

Out[142]: (551, 15)

```
In [143]: # Selecting a DataFrame slice
table
```

Out[143]:


driver_gender	F	M
driver_race		
Asian	551	1838
Black	2681	9604
Hispanic	1953	7774
Other	53	212
White	18536	43334

```
In [144]: table = table.loc['Asian':'Hispanic']
table
```


Out[144]:

driver_gender	F	M
driver_race		
Asian	551	1838
Black	2681	9604
Hispanic	1953	7774


```
In [145]: # Creating a line plot
table.plot()
plt.show()
```



```
In [146]: # Creating a bar plot
table.plot(kind='bar')
plt.show()
```



```
In [147]: # Stacking the bars
table.plot(kind='bar', stacked=True)
plt.show()
```



Exercises

Tallying violations by district

```
In [157]: # Create a frequency table of districts and violations
print(pd.crosstab(ri.district,ri.violation))

# Save the frequency table as 'all_zones'
all_zones = pd.crosstab(ri.district,ri.violation)
```

violation	Equipment	Moving violation	Other	Registration/plates	Seat belt	\
district						
Zone K1	672	1254	290	120	0	
Zone K2	2061	2962	942	768	481	
Zone K3	2302	2898	705	695	638	
Zone X1	296	671	143	38	74	
Zone X3	2049	3086	769	671	820	
Zone X4	3541	5353	1560	1411	843	

violation	Speeding
district	
Zone K1	5960
Zone K2	10448
Zone K3	12322
Zone X1	1119
Zone X3	8779
Zone X4	9795

```
In [158]: # Select rows 'Zone K1' through 'Zone K3'
print(all_zones.loc['Zone K1':'Zone K3'])

# Save the smaller table as 'k_zones'
k_zones = all_zones.loc['Zone K1':'Zone K3']
```

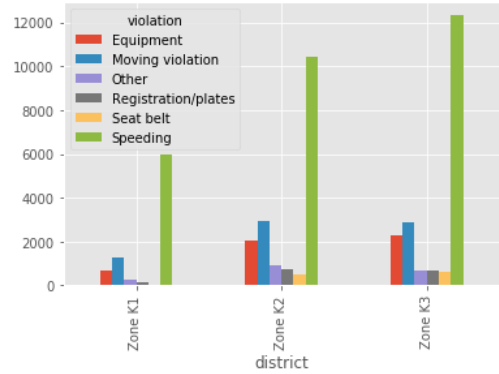
violation	Equipment	Moving violation	Other	Registration/plates	Seat belt	\
district						
Zone K1	672	1254	290	120	0	
Zone K2	2061	2962	942	768	481	
Zone K3	2302	2898	705	695	638	

violation	Speeding
district	
Zone K1	5960
Zone K2	10448
Zone K3	12322

Plotting violations by district

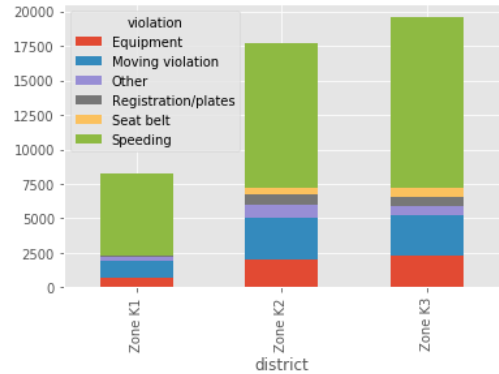
```
In [166]: # Create a bar plot of 'k_zones'
k_zones.plot(kind='bar')

# Display the plot
plt.show()
```



```
In [167]: # Create a stacked bar plot of 'k_zones'
k_zones.plot(kind='bar',stacked=True)

# Display the plot
plt.show()
```



How long might you be stopped for a violation?

```
In [177]: # Analyzing an object column
apple['change'] = (apple.price.diff()>=0).map({True:'up',False:'down'})
apple.head(5)
```

Out[177]:

	price	volume	change
date_and_time			
2018-02-01 16:00:00	170.160004	25555900	down
2018-03-01 16:00:00	172.529999	29517900	up
2018-04-01 16:00:00	172.539993	22434600	up
2018-05-01 16:00:00	173.440002	23660000	up
2018-08-01 16:00:00	174.350006	20567800	up

```
In [178]: apple.change.dtype
```

Out[178]: dtype('O')

```
In [179]: # Mapping one set of values to another
mapping = {'up':True, 'down':False}
apple['is_up'] = apple.change.map(mapping)
apple.head(5)
```

Out[179]:

	price	volume	change	is_up
date_and_time				
2018-02-01 16:00:00	170.160004	25555900	down	False
2018-03-01 16:00:00	172.529999	29517900	up	True
2018-04-01 16:00:00	172.539993	22434600	up	True
2018-05-01 16:00:00	173.440002	23660000	up	True
2018-08-01 16:00:00	174.350006	20567800	up	True

```
In [180]: apple.is_up.mean()
```

Out[180]: 0.5245901639344263

Calculating the search rate


In [182]: `# Calculating the search rate`
`search_rate = ri.groupby('violation').search_conducted.mean()`
`search_rate`

Out[182]:

violation	
Equipment	0.064280
Moving violation	0.057014
Other	0.045362
Registration/plates	0.093438
Seat belt	0.031513
Speeding	0.021560

Name: search_conducted, dtype: float64

In [183]: `# Creating a bar plot`
`search_rate.plot(kind='bar')`
`plt.show()`



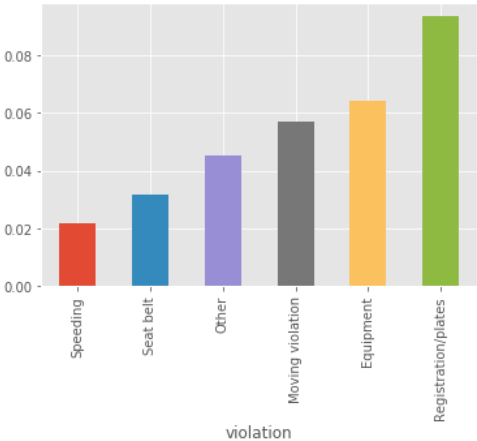
In [184]: `# Ordering the bars (1)`
`search_rate.sort_values()`

Out[184]:


violation	
Speeding	0.021560
Seat belt	0.031513
Other	0.045362
Moving violation	0.057014
Equipment	0.064280
Registration/plates	0.093438

Name: search_conducted, dtype: float64

In [185]: `# Ordering the bars (2)`
`search_rate.sort_values().plot(kind='bar')`
`plt.show()`



In [186]: `# Rotating the bars`
`search_rate.sort_values().plot(kind='barh')`
`plt.show()`



Exercises

Converting stop durations to numbers

In [189]: `# Print the unique values in 'stop_duration'`
`print(ri.stop_duration.unique())`

`# Create a dictionary that maps strings to integers`
`mapping = {'0-15 Min':8, '16-30 Min':23, '30+ Min':45}`

`# Convert the 'stop_duration' strings to integers using the 'mapping'`
`ri['stop_minutes'] = ri.stop_duration.map(mapping)`

`# Print the unique values in 'stop_minutes'`
`print(ri.stop_minutes.unique())`

`['0-15 Min' '16-30 Min' '30+ Min']`
`[8 23 45]`

Plotting stop length

```
In [191]: # Calculate the mean 'stop_minutes' for each value in 'violation_raw'
print(ri.stop_minutes.groupby(ri.violation_raw).mean())

# Save the resulting Series as 'stop_length'
stop_length = ri.stop_minutes.groupby(ri.violation_raw).mean()


# Sort 'stop_length' by its values and create a horizontal bar plot
stop_length.sort_values()

# Display the plot
stop_length.sort_values().plot(kind='barh')
plt.show()
```

violation_raw

APB	17.967033
Call for Service	22.124371
Equipment/Inspection Violation	11.445655
Motorist Assist/Courtesy	17.741463
Other Traffic Violation	13.844490
Registration Violation	13.736970
Seatbelt Violation	9.662815
Special Detail/Directed Patrol	15.123632
Speeding	10.581562
Suspicious Person	14.910714
Violation of City/Town Ordinance	13.254144
Warrant	24.055556

Name: stop_minutes, dtype: float64



Chap 4: Analyzing the effect of weather on policing

```
In [192]: # Import plotting modules
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
plt.style.use('ggplot')
```

Exploring the weather dataset

```
In [194]: # Examining the columns
weather = pd.read_csv('datasets/weather.csv')
weather.head()
```

Out[194]:

	STATION	DATE	TAVG	TMIN	TMAX	AWND	WSF2	WT01	WT02	WT03	...	WT11	WT13	WT14	WT15	WT16	WT17	WT18	WT19	WT21	WT22
0	USW00014765	2005-01-01	44.0	35	53	8.95	25.1	1.0	NaN	NaN	...	NaN	1.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1	USW00014765	2005-01-02	36.0	28	44	9.40	14.1	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	1.0	NaN	1.0	NaN	NaN	NaN
2	USW00014765	2005-01-03	49.0	44	53	6.93	17.0	1.0	NaN	NaN	...	NaN	1.0	NaN	NaN	1.0	NaN	NaN	NaN	NaN	NaN
3	USW00014765	2005-01-04	42.0	39	45	6.93	16.1	1.0	NaN	NaN	...	NaN	1.0	1.0	NaN	1.0	NaN	NaN	NaN	NaN	NaN
4	USW00014765	2005-01-05	36.0	28	43	7.83	17.0	1.0	NaN	NaN	...	NaN	1.0	NaN	NaN	1.0	NaN	1.0	NaN	NaN	NaN

5 rows × 27 columns

TAVG, TMIN, TMAX: Temperature
AWND, WSF2: Wind speed
WT01 ... WT22: Bad weather conditions

```
In [195]: # Examining wind speed
weather[['AWND', 'WSF2']].head()
```

Out[195]:


	AWND	WSF2
0	8.95	25.1
1	9.40	14.1
2	6.93	17.0
3	6.93	16.1
4	7.83	17.0

```
In [196]: weather[['AWND', 'WSF2']].describe()
```


Out[196]:

	AWND	WSF2
count	4017.000000	4017.000000
mean	8.593707	19.274782
std	3.364601	5.623866
min	0.220000	4.900000
25%	6.260000	15.000000
50%	8.050000	17.900000
75%	10.290000	21.900000
max	26.840000	48.100000


```
In [197]: # Creating a box plot
weather[['AWND', 'WSF2']].plot(kind='box')
plt.show()
```



```
In [198]: # Creating a histogram (1)
weather['WDIFF'] = weather.WSF2 - weather.AWND
weather.WDIFF.plot(kind='hist')
plt.show()
```



```
In [199]: # Creating a histogram (2)
weather.WDIFF.plot(kind='hist', bins=20)
plt.show()
```



Exercises

Plotting the temperature


```
In [205]: # Read 'weather.csv' into a DataFrame named 'weather'
weather = pd.read_csv('datasets/weather.csv')

# Describe the temperature columns
print(weather[['TMIN', 'TAVG', 'TMAX']].describe())

# Create a box plot of the temperature columns
weather[['TMIN', 'TAVG', 'TMAX']].plot(kind='box')

# Display the plot
plt.show()
```

	TMIN	TAVG	TMAX
count	4017.000000	1217.000000	4017.000000
mean	43.484441	52.493016	61.268608
std	17.020298	17.830714	18.199517
min	-5.000000	6.000000	15.000000
25%	30.000000	39.000000	47.000000
50%	44.000000	54.000000	62.000000
75%	58.000000	68.000000	77.000000
max	77.000000	86.000000	102.000000



Plotting the temperature difference

```
In [211]: # Create a 'TDIFF' column that represents temperature difference
weather['TDIFF'] = weather.TMAX - weather.TMIN


# Describe the 'TDIFF' column
print(weather.TDIFF.describe())

# Create a histogram with 20 bins to visualize 'TDIFF'
weather.TDIFF.plot(kind='hist',bins=20)

# Display the plot
plt.show()
```

count	4017.000000
mean	17.784167
std	6.350720
min	2.000000
25%	14.000000
50%	18.000000
75%	22.000000
max	43.000000

Name: TDIFF, dtype: float64



Categorizing the weather

```
In [214]: # Selecting a DataFrame slice
weather.shape
```

Out[214]: (4017, 28)

```
In [215]: weather.columns
```

Out[215]: Index(['STATION', 'DATE', 'TAVG', 'TMIN', 'TMAX', 'AWND', 'WSF2', 'WT01', 'WT02', 'WT03', 'WT04', 'WT05', 'WT06', 'WT07', 'WT08', 'WT09', 'WT10', 'WT11', 'WT13', 'WT14', 'WT15', 'WT16', 'WT17', 'WT18', 'WT19', 'WT21', 'WT22', 'TDIFF'], dtype='object')


```
In [216]: temp = weather.loc[:, 'TAVG':'TMAX']
temp.shape

Out[216]: (4017, 3)

In [217]: temp.columns

Out[217]: Index(['TAVG', 'TMIN', 'TMAX'], dtype='object')

In [218]: # DataFrame operations
temp.head()

Out[218]:
   TAVG  TMIN  TMAX
0   44.0    35    53
1   36.0    28    44
2   49.0    44    53
3   42.0    39    45
4   36.0    28    43

In [219]: temp.sum()

Out[219]: TAVG      63884.0
TMIN      174677.0
TMAX      246116.0
dtype: float64

In [220]: temp.sum(axis='columns').head()

Out[220]: 0      132.0
1      108.0
2      146.0
3      126.0
4      107.0
dtype: float64
```

Mapping one set of values to another

```
In [221]: # Mapping one set of values to another
ri.stop_duration.unique()

Out[221]: array(['0-15 Min', '16-30 Min', '30+ Min'], dtype=object)

In [222]: mapping = {'0-15 Min':'short', '16-30 Min':'medium', '30+ Min':'long'}
ri['stop_length'] = ri.stop_duration.map(mapping)
ri.stop_length.dtype

Out[222]: dtype('O')

In [223]: # Changing data type from object to category
ri.stop_length.unique()

Out[223]: array(['short', 'medium', 'long'], dtype=object)

In [224]: ri.stop_length.memory_usage(deep=True)

Out[224]: 8689481

In [258]: cats = ['short', 'medium', 'long']
ri['stop_length'] = ri.stop_length.astype(pd.api.types.CategoricalDtype(categories=cats))
```

Category type stores the data more efficiently Allows you to specify a logical order for the categories

```
In [259]: ri.stop_length.memory_usage(deep=True)

Out[259]: 3400530

In [260]: # Using ordered categories
ri.stop_length.head()

Out[260]: stop_datetime
2005-01-04 12:55:00    short
2005-01-23 23:15:00    short
2005-02-17 04:15:00    short
2005-02-20 17:15:00    medium
2005-02-24 01:20:00    short
Name: stop_length, dtype: category
Categories (3, object): [short < medium < long]

In [261]: ri[ri.stop_length > 'short'].shape

Out[261]: (16959, 17)

In [262]: ri.groupby('stop_length').is_arrested.mean()

Out[262]: stop_length
short      0.013654
medium     0.093595
long       0.261572
Name: is_arrested, dtype: float64
```

Exercises

Counting bad weather conditions

- WT05 indicates "Hail"
- WT11 indicates "High or damaging winds"
- WT17 indicates "Freezing rain"


```
In [263]: # Copy 'WT01' through 'WT22' to a new DataFrame
WT = weather.loc[:, 'WT01':'WT22']

# Calculate the sum of each row in 'WT'
weather['bad_conditions'] = WT.sum(axis='columns')

# Replace missing values in 'bad_conditions' with '0'
weather['bad_conditions'] = weather.bad_conditions.fillna(0).astype('int')

# Create a histogram to visualize 'bad_conditions'
weather.bad_conditions.plot(kind='hist')

# Display the plot
plt.show()
```



Rating the weather conditions

```
In [264]: # Count the unique values in 'bad_conditions' and sort the index
print(weather.bad_conditions.value_counts().sort_index())

# Create a dictionary that maps integers to strings
mapping = {0:'good', 1:'bad', 2:'bad', 3:'bad', 4:'bad', 5:'worse', 6:'worse', 7:'worse', 8:'worse', 9:'worse'}

# Convert the 'bad_conditions' integers to strings using the 'mapping'
weather['rating'] = weather.bad_conditions.map(mapping)

# Count the unique values in 'rating'
print(weather.rating.value_counts().sort_index())

0      1749
1       613
2       367
3       380
4       476
5       282
6       101
7        41
8         4
9         4
Name: bad_conditions, dtype: int64
bad      1836
good     1749
worse     432
Name: rating, dtype: int64
```

Changing the data type to category

```
In [265]: # Create a List of weather ratings in Logical order
cats = ['good', 'bad', 'worse']

# Change the data type of 'rating' to category
weather['rating'] = weather.rating.astype(pd.api.types.CategoricalDtype(categories=cats))

# Examine the head of 'rating'
print(weather.rating.head())

0      bad
1      bad
2      bad
3      bad
4      bad
Name: rating, dtype: category
Categories (3, object): [good, bad, worse]
```

Merging datasets

```
In [411]: # APPLE DATASETS
apple = pd.read_csv('datasets/aapl_ohlcv.csv',usecols=[0,1,2])
apple['date_and_time'] = pd.to_datetime(apple.date.str.replace('/', '-').str.cat(apple.time,sep=' '))
apple['date'] = pd.to_datetime(apple.date)
apple.set_index('date_and_time',inplace=True)

high_low = pd.read_csv('datasets/AAPL2018.csv',usecols=[0,2,3],skiprows=1,names=['DATE','HIGH','LOW'])
high_low['DATE'] = pd.to_datetime(high_low.DATE)
```

```
In [412]: # Preparing the first DataFrame
apple.head()
```

Out[412]:

	date	time	price
date_and_time			
2018-02-01 16:00:00	2018-02-01	16:00	170.160004
2018-03-01 16:00:00	2018-03-01	16:00	172.529999
2018-04-01 16:00:00	2018-04-01	16:00	172.539993
2018-05-01 16:00:00	2018-05-01	16:00	173.440002
2018-08-01 16:00:00	2018-08-01	16:00	174.350006

```
In [413]: apple.reset_index(inplace=True)
apple.head()
```

Out[413]:

	date_and_time	date	time	price
0	2018-02-01 16:00:00	2018-02-01	16:00	170.160004
1	2018-03-01 16:00:00	2018-03-01	16:00	172.529999
2	2018-04-01 16:00:00	2018-04-01	16:00	172.539993
3	2018-05-01 16:00:00	2018-05-01	16:00	173.440002
4	2018-08-01 16:00:00	2018-08-01	16:00	174.350006

```
In [414]: # Preparing the second DataFrame
high_low.head()
```

Out[414]:

	DATE	HIGH	LOW
0	2018-02-01	172.300003	169.259995
1	2018-03-01	174.550003	171.960007
2	2018-04-01	173.470001	172.080002
3	2018-05-01	175.369995	173.050003
4	2018-08-01	175.610001	173.929993

```
In [415]: high = high_low[['DATE', 'HIGH']]
high.head()
```

Out[415]:

	DATE	HIGH
0	2018-02-01	172.300003
1	2018-03-01	174.550003
2	2018-04-01	173.470001
3	2018-05-01	175.369995
4	2018-08-01	175.610001

```
In [416]: # Merging the DataFrames
apple_high = pd.merge(left=apple, right=high, left_on='date', right_on='DATE', how='left')
```

```
In [417]: # Comparing the DataFrames
apple_high.head()
```

Out[417]:

	date_and_time	date	time	price	DATE	HIGH
0	2018-02-01 16:00:00	2018-02-01	16:00	170.160004	2018-02-01	172.300003
1	2018-03-01 16:00:00	2018-03-01	16:00	172.529999	2018-03-01	174.550003
2	2018-04-01 16:00:00	2018-04-01	16:00	172.539993	2018-04-01	173.470001
3	2018-05-01 16:00:00	2018-05-01	16:00	173.440002	2018-05-01	175.369995
4	2018-08-01 16:00:00	2018-08-01	16:00	174.350006	2018-08-01	175.610001

```
In [418]: apple.head()
```

Out[418]:

	date_and_time	date	time	price
0	2018-02-01 16:00:00	2018-02-01	16:00	170.160004
1	2018-03-01 16:00:00	2018-03-01	16:00	172.529999
2	2018-04-01 16:00:00	2018-04-01	16:00	172.539993
3	2018-05-01 16:00:00	2018-05-01	16:00	173.440002
4	2018-08-01 16:00:00	2018-08-01	16:00	174.350006

```
In [419]: high.head()
```

Out[419]:

	DATE	HIGH
0	2018-02-01	172.300003
1	2018-03-01	174.550003
2	2018-04-01	173.470001
3	2018-05-01	175.369995
4	2018-08-01	175.610001

```
In [420]: # Setting the index
apple_high.set_index('date_and_time', inplace=True)
apple_high.head()
```

Out[420]:

	date_and_time	date	time	price	DATE	HIGH
	2018-02-01 16:00:00	2018-02-01	16:00	170.160004	2018-02-01	172.300003
	2018-03-01 16:00:00	2018-03-01	16:00	172.529999	2018-03-01	174.550003
	2018-04-01 16:00:00	2018-04-01	16:00	172.539993	2018-04-01	173.470001
	2018-05-01 16:00:00	2018-05-01	16:00	173.440002	2018-05-01	175.369995
	2018-08-01 16:00:00	2018-08-01	16:00	174.350006	2018-08-01	175.610001

Exercises

Preparing the DataFrames

```
In [422]: # Reset the index of 'ri'
ri.reset_index(inplace=True)

# Examine the head of 'ri'
print(ri.head())

# Create a DataFrame from the 'DATE' and 'rating' columns
weather_rating = weather[['DATE', 'rating']]

# Examine the head of 'weather_rating'
print(weather_rating.head())
```

	stop_datetime	stop_date	stop_time	driver_gender	driver_race	\
0	2005-01-04 12:55:00	2005-01-04	12:55	M	White	
1	2005-01-23 23:15:00	2005-01-23	23:15	M	White	
2	2005-02-17 04:15:00	2005-02-17	04:15	M	White	
3	2005-02-20 17:15:00	2005-02-20	17:15	M	White	
4	2005-02-24 01:20:00	2005-02-24	01:20	F	White	

	violation_raw	violation	search_conducted	search_type	\
0	Equipment/Inspection Violation	Equipment	False	NaN	
1	Speeding	Speeding	False	NaN	
2	Speeding	Speeding	False	NaN	
3	Call for Service	Other	False	NaN	
4	Speeding	Speeding	False	NaN	

	stop_outcome	is_arrested	stop_duration	drugs_related_stop	district	\
0	Citation	False	0-15 Min	False	Zone X4	
1	Citation	False	0-15 Min	False	Zone K3	
2	Citation	False	0-15 Min	False	Zone X4	
3	Arrest Driver	True	16-30 Min	False	Zone X1	
4	Citation	False	0-15 Min	False	Zone X3	

	inventory	frisk	stop_minutes	stop_length
0	False	False	8	short
1	False	False	8	short
2	False	False	8	short
3	False	False	23	medium
4	False	False	8	short

	DATE	rating
0	2005-01-01	bad
1	2005-01-02	bad
2	2005-01-03	bad
3	2005-01-04	bad
4	2005-01-05	bad

Merging the DataFrame

```
In [423]: # Examine the shape of 'ri'
print(ri.shape)

# Merge 'ri' and 'weather_rating' using a Left join
ri_weather = pd.merge(left=ri, right=weather_rating, left_on='stop_date', right_on='DATE', how='left')

# Examine the shape of 'ri_weather'
print(ri_weather.shape)

# Set 'stop_datetime' as the index of 'ri_weather'
ri_weather.set_index('stop_datetime', inplace=True)

(86536, 18)
(86536, 20)
```

Does weather affect the arrest rate?

```
In [425]: # Driver gender and vehicle searches
ri.search_conducted.mean()
```

Out[425]: 0.0382153092354627

```
In [426]: ri.groupby('driver_gender').search_conducted.mean()
```

Out[426]: driver_gender
F 0.019181
M 0.045426
Name: search_conducted, dtype: float64

```
In [427]: ri.groupby(['violation', 'driver_gender']).search_conducted.mean()

Out[427]: violation      driver_gender
Equipment      F      0.039984
              M      0.071496
Moving violation F      0.039257
              M      0.061524
Other           F      0.041018
              M      0.046191
Registration/plates F      0.054924
              M      0.108802
Seat belt       F      0.017301
              M      0.035119
Speeding        F      0.008309
              M      0.027885
Name: search_conducted, dtype: float64

In [429]: # Examining a multi-indexed Series
search_rate = ri.groupby(['violation', 'driver_gender']).search_conducted.mean()
search_rate

Out[429]: violation      driver_gender
Equipment      F      0.039984
              M      0.071496
Moving violation F      0.039257
              M      0.061524
Other           F      0.041018
              M      0.046191
Registration/plates F      0.054924
              M      0.108802
Seat belt       F      0.017301
              M      0.035119
Speeding        F      0.008309
              M      0.027885
Name: search_conducted, dtype: float64
```

```
In [430]: type(search_rate)

Out[430]: pandas.core.series.Series
```

```
In [431]: type(search_rate.index)

Out[431]: pandas.core.indexes.multi.MultiIndex
```

```
In [432]: # Working with a multi-indexed Series
search_rate

Out[432]: violation      driver_gender
Equipment      F      0.039984
              M      0.071496
Moving violation F      0.039257
              M      0.061524
Other           F      0.041018
              M      0.046191
Registration/plates F      0.054924
              M      0.108802
Seat belt       F      0.017301
              M      0.035119
Speeding        F      0.008309
              M      0.027885
Name: search_conducted, dtype: float64
```

```
In [433]: search_rate.loc['Equipment']

Out[433]: driver_gender
F      0.039984
M      0.071496
Name: search_conducted, dtype: float64
```

```
In [434]: search_rate.loc['Equipment', 'M']

Out[434]: 0.07149643705463182
```

```
In [435]: # Converting a multi-indexed Series to a DataFrame
search_rate.unstack()

Out[435]:
```

driver_gender	F	M
violation		
Equipment	0.039984	0.071496
Moving violation	0.039257	0.061524
Other	0.041018	0.046191
Registration/plates	0.054924	0.108802
Seat belt	0.017301	0.035119
Speeding	0.008309	0.027885

```
In [436]: type(search_rate.unstack())

Out[436]: pandas.core.frame.DataFrame
```

```
In [437]: ri.pivot_table(index='violation', columns='driver_gender', values='search_conducted')

Out[437]:
```

driver_gender	F	M
violation		
Equipment	0.039984	0.071496
Moving violation	0.039257	0.061524
Other	0.041018	0.046191
Registration/plates	0.054924	0.108802
Seat belt	0.017301	0.035119
Speeding	0.008309	0.027885

Exercises

Comparing arrest rates by weather rating

```
In [439]: # Calculate the overall arrest rate
print(ri_weather.is_arrested.mean())

0.0355690117407784

In [440]: # Calculate the arrest rate for each 'rating'
print(ri_weather.is_arrested.groupby(ri_weather.rating).mean())

rating
good      0.033715
bad       0.036261
worse     0.041667
Name: is_arrested, dtype: float64
```

In [441]:

```
# Calculate the arrest rate for each 'violation' and 'rating'
print(ri_weather.is_arrested.groupby([ri_weather.violation,ri_weather.rating]).mean())
```

violation	rating	
Equipment	good	0.059007
	bad	0.066311
	worse	0.097357
Moving violation	good	0.056227
	bad	0.058050
	worse	0.065860
Other	good	0.076966
	bad	0.087443
	worse	0.062893
Registration/plates	good	0.081574
	bad	0.098160
	worse	0.115625
Seat belt	good	0.028587
	bad	0.022493
	worse	0.000000
Speeding	good	0.013405
	bad	0.013314
	worse	0.016886

Name: is_arrested, dtype: float64

Selecting from a multi-indexed Series

In [442]:

```
# Save the output of the groupby operation from the last exercise
arrest_rate = ri_weather.groupby(['violation', 'rating']).is_arrested.mean()

# Print the 'arrest_rate' Series
print(arrest_rate)

# Print the arrest rate for moving violations in bad weather
print(arrest_rate.loc['Moving violation', 'bad'])

# Print the arrest rates for speeding violations in all three weather conditions
print(arrest_rate.loc['Speeding'])
```

violation	rating	
Equipment	good	0.059007
	bad	0.066311
	worse	0.097357
Moving violation	good	0.056227
	bad	0.058050
	worse	0.065860
Other	good	0.076966
	bad	0.087443
	worse	0.062893
Registration/plates	good	0.081574
	bad	0.098160
	worse	0.115625
Seat belt	good	0.028587
	bad	0.022493
	worse	0.000000
Speeding	good	0.013405
	bad	0.013314
	worse	0.016886

Name: is_arrested, dtype: float64

0.05804964058049641

rating

good	0.013405
bad	0.013314
worse	0.016886

Name: is_arrested, dtype: float64

Reshaping the arrest rate data

In [443]:

```
# Unstack the 'arrest_rate' Series into a DataFrame
print(arrest_rate.unstack())

# Create the same DataFrame using a pivot table
print(ri_weather.pivot_table(index='violation', columns='rating', values='is_arrested'))
```

rating	good	bad	worse
violation			
Equipment	0.059007	0.066311	0.097357
Moving violation	0.056227	0.058050	0.065860
Other	0.076966	0.087443	0.062893
Registration/plates	0.081574	0.098160	0.115625
Seat belt	0.028587	0.022493	0.000000
Speeding	0.013405	0.013314	0.016886

rating	good	bad	worse
violation			
Equipment	0.059007	0.066311	0.097357
Moving violation	0.056227	0.058050	0.065860
Other	0.076966	0.087443	0.062893
Registration/plates	0.081574	0.098160	0.115625
Seat belt	0.028587	0.022493	0.000000
Speeding	0.013405	0.013314	0.016886