

DataCamp Data Science Courses

Extreme Gradient Boosting with XGBoost

Chap 1: Classification with XGBoost

```
In [1]: # Import plotting modules
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
plt.style.use('ggplot')

import xgboost as xgb
from sklearn.model_selection import train_test_split
```

Welcome to the course!

Supervised Learning

- Regression: covered in chapter 2
- Classification
 - binary: predict a person will buy insurance
 - metric: AUC (Area under the ROC Curve)
 - multi-class: classifying the species of a given bird in an image
 - metric: accuracy score, confusion matrix

common classification models: logistic regression, decision trees.

Supervised Learning considerations

- Features can be either numeric or categorical
- Numeric features should be scaled (Z-scored)
 - e.g. essential to train SVM models
- Categorical features should be encoded (one-hot)

Other kinds of Supervised Learning problems

- Ranking: Predicting an ordering on a set of choices
 - Google search suggestions
- Recommendation: Recommending an item to a user based on his/her consumption history and profile
 - Netflix

Introducing XGBoost

- Optimized gradient-boosting machine learning library
- Originally written in C++
- Why it is popular
 - Speed and performance
 - Core algorithm is parallelizable
 - Consistently outperforms single-algorithm methods
 - State-of-the-art performance in many ML benchmark datasets

```
In [ ]: # Using XGBoost: A Quick Example

class_data = pd.read_csv("datasets/classification_data.csv")

X, y = class_data.iloc[:, :-1], class_data.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=123)

xg_cl = xgb.XGBClassifier(objective='binary:logistic', n_estimators=10, seed=123)

xg_cl.fit(X_train, y_train)
preds = xg_cl.predict(X_test)

accuracy = float(np.sum(preds==y_test))/y_test.shape[0]
print("accuracy: %f" % (accuracy))
```

```
In [ ]: # EXERCISES
```

```
In [ ]: # XGBoost: Fit/Predict

# Import xgboost
import xgboost as xgb

# Create arrays for the features and the target: X, y
X, y = churn_data.iloc[:, :-1], churn_data.iloc[:, -1]

# Create the training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=123)

# Instantiate the XGBClassifier: xg_cl
xg_cl = xgb.XGBClassifier(objective='binary:logistic', n_estimators=10, seed=123)

# Fit the classifier to the training set
xg_cl.fit(X_train, y_train)

# Predict the labels of the test set: preds
preds = xg_cl.predict(X_test)

# Compute the accuracy: accuracy
accuracy = float(np.sum(preds == y_test)) / y_test.shape[0]
print("accuracy: %f" % (accuracy))
```

What is a decision tree?

Decision Trees

- Base learner - Individual learning algorithm in an ensemble algorithm
- Composed of a series of binary questions
- Predictions happen at the "leaves" of the tree
 - leaf nodes always contain decision values
- Constructed iteratively (one decision at a time)
 - Until a stopping criterion is met
- Individual decision trees tend to overfit
 - low bias, high variance
 - tend to overfit training data, and generalize poorly to new data

XGBoost

- Uses classification and regression trees (CART)
- Contain real-valued score in each leaf
 - regardless of classification or regression problem
 - can be thresholded to convert into categories for classification problems

```
In [4]: # EXERCISES
```

```
In [5]: # Decision trees

# Import the necessary modules
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn import datasets

bc = datasets.load_breast_cancer()
X = bc.data
y = bc.target

X.shape, y.shape
```

```
Out[5]: ((569, 30), (569,))
```

```
In [6]: # Create the training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=123)

# Instantiate the classifier: dt_clf_4
dt_clf_4 = DecisionTreeClassifier(max_depth=4)

# Fit the classifier to the training set
dt_clf_4.fit(X_train, y_train)

# Predict the labels of the test set: y_pred_4
y_pred_4 = dt_clf_4.predict(X_test)

# Compute the accuracy of the predictions: accuracy
accuracy = float(np.sum(y_pred_4 == y_test)) / y_test.shape[0]
print("accuracy:", accuracy)
```

```
accuracy: 0.9736842105263158
```

What is Boosting?

Boosting overview

- Not a specific machine learning algorithm
- Concept that can be applied to a set of machine learning models
 - "Meta-algorithm"
- Ensemble meta-algorithm used to convert many weak learners into a strong learner

- Weak learner: ML algorithm that is slightly better than chance
 - Example: Decision tree whose predictions are slightly better than 50%
- Boosting converts a collection of weak learners into a strong learner
- Strong learner: Any algorithm that can be tuned to achieve good performance

How boosting is accomplished

- Iteratively learning a set of weak models on subsets of the data
- Weighing each weak prediction according to each weak learner's performance
- Combine the weighted predictions to obtain a single weighted prediction that is much better than the individual predictions themselves!

Model evaluation through cross-validation

- Cross-validation: Robust method for estimating the performance of a model on unseen data
- Generates many non-overlapping train/test splits on training data
- Reports the average test set performance across all data splits

```
In [ ]: # Cross-validation in XGBoost example

class_data = pd.read_csv("classification_data.csv")

churn_dmatrix = xgb.DMatrix(data=churn_data.iloc[:, :-1],
                             label=churn_data.month_5_still_here)

params={"objective":"binary:logistic","max_depth":4}
cv_results = xgb.cv(dtrain=churn_dmatrix, params=params, nfold=4,
                    num_boost_round=10, metrics="error", as_pandas=True)

print("Accuracy: %f" %((1-cv_results["test-error-mean"]).iloc[-1]))
```

num_boost_round: number of trees to run

```
In [8]: # EXERCISES
```

```
In [9]: # Measuring accuracy

# Create the DMatrix: churn_dmatrix
churn_dmatrix = xgb.DMatrix(data=X, label=y)

# Create the parameter dictionary: params
params = {"objective":"reg:logistic", "max_depth":3}

# Perform cross-validation: cv_results
cv_results = xgb.cv(dtrain=churn_dmatrix, params=params, nfold=3, num_boost_round=5, metrics="error", as_pandas=True, seed=123)

# Print cv_results
print(cv_results)

# Print the accuracy
print(((1-cv_results["test-error-mean"]).iloc[-1]))
```

| | test-error-mean | test-error-std | train-error-mean | train-error-std |
|---|-----------------|----------------|------------------|-----------------|
| 0 | 0.066824 | 0.019564 | 0.025480 | 0.002451 |
| 1 | 0.061524 | 0.013876 | 0.021969 | 0.001257 |
| 2 | 0.056252 | 0.010004 | 0.014945 | 0.006589 |
| 3 | 0.052734 | 0.011418 | 0.012306 | 0.003300 |
| 4 | 0.054497 | 0.012485 | 0.010549 | 0.004314 |
| | 0.945502666667 | | | |

```
In [10]: # Measuring AUC

# Perform cross_validation: cv_results
cv_results = xgb.cv(dtrain=churn_dmatrix, params=params, nfold=3, num_boost_round=5, metrics="auc", as_pandas=True, seed=123)

# Print cv_results
print(cv_results)

# Print the AUC
print((cv_results["test-auc-mean"]).iloc[-1])
```

| | test-auc-mean | test-auc-std | train-auc-mean | train-auc-std |
|---|---------------|--------------|----------------|---------------|
| 0 | 0.961473 | 0.024760 | 0.987225 | 0.001301 |
| 1 | 0.969078 | 0.022616 | 0.993244 | 0.004295 |
| 2 | 0.972491 | 0.024377 | 0.995224 | 0.003751 |
| 3 | 0.971354 | 0.025405 | 0.997125 | 0.002042 |
| 4 | 0.974002 | 0.026527 | 0.997610 | 0.001871 |
| | 0.974002 | | | |

When should I use XGBoost?

When to use XGBoost

- You have a large number of training samples
 - Greater than 1000 training samples and less 100 features
 - The number of features < number of training samples
- You have a mixture of categorical and numeric features
 - Or just numeric features

When NOT to use XGBoost

- Image recognition
- Computer vision
- Natural language processing and understanding problems
- When the number of training samples is significantly smaller than the number of features

Chap 2: Regression with XGBoost

```
In [11]: # Import plotting modules
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
plt.style.use('ggplot')

from sklearn.model_selection import train_test_split
import xgboost as xgb
from sklearn.metrics import mean_squared_error
```

Regression review

Common regression metrics

- Root mean squared error (RMSE)
 - Treats negative and positive differences equally
 - Tends to punish larger differences more than smaller ones
- Mean absolute error (MAE)
 - Not affected by large error values as RMSE
 - Lacks nice mathematical properties
 - Less frequently used as evaluation metric

Common regression algorithms

- Linear regression
- Decision trees
 - Can be used for both regression/classification tasks
 - Important building block for XGBoost models

Objective (loss) functions and base learners

Loss/Objective Functions

- Quantifies how far off a prediction is from the actual result
- Measures the difference between estimated and true values for some collection of data
- **Goal:** Find the model that yields the minimum value of the loss function

- Loss function names in xgboost:
 - **reg:linear** - use for regression problems
 - **reg:logistic** - use for classification problems when you want just decision, not probability
 - **binary:logistic** - use when you want probability rather than just decision

Base Learners

- XGBoost involves creating a meta-model that is composed of many individual models that combine to give a final prediction
- Individual models = base learners
 - Want base learners that are slightly better than random guessing for some part of data but are uniformly bad for the remaining majority of the data.
 - When combined they create final prediction that is **non-linear**
 - Each base learner should be good at distinguishing or predicting different parts of the dataset
- Two kinds of base learners: **tree** and **linear**

```
In [12]: # TREES BASE LEARNERS example: Scikit-Learn API

boston_data = pd.read_csv("datasets/boston.csv")
X, y = boston_data.iloc[:, :-1], boston_data.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=123)

xg_reg = xgb.XGBRegressor(objective='reg:linear', n_estimators=10, seed=123)

xg_reg.fit(X_train, y_train)
preds = xg_reg.predict(X_test)
```

```
In [13]: rmse = np.sqrt(mean_squared_error(y_test, preds))
print("RMSE: %f" % (rmse))
```

RMSE: 9.749041

```
In [14]: # LINEAR BASE LEARNERS Example: Learning API Only

boston_data = pd.read_csv("datasets/boston.csv")
X, y = boston_data.iloc[:, :-1], boston_data.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=123)

DM_train = xgb.DMatrix(data=X_train, label=y_train)
DM_test = xgb.DMatrix(data=X_test, label=y_test)

# OPTIONS FOR BOOSTER: gbtrees(default), gblines or dart
params = {"booster": "gblines", "objective": "reg:linear"}

xg_reg = xgb.train(params = params, dtrain=DM_train, num_boost_round=10)
preds = xg_reg.predict(DM_test)
```

```
In [15]: rmse = np.sqrt(mean_squared_error(y_test, preds))
print("RMSE: %f" % (rmse))
```

RMSE: 5.492696

```
In [16]: # EXERCISES
```

```
In [17]: # Decision trees as base learners

df = pd.read_csv('datasets/ames_housing_trimmed_processed.csv')
X, y = df.iloc[:, :-1], df.iloc[:, -1]
X.shape, y.shape
```

Out[17]: ((1460, 56), (1460,))

```
In [18]: # Create the training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=123)

# Instantiate the XGBRegressor: xg_reg
xg_reg = xgb.XGBRegressor(objective='reg:linear', n_estimators=10)

# Fit the regressor to the training set
xg_reg.fit(X_train, y_train)

# Predict the labels of the test set: preds
preds = xg_reg.predict(X_test)

# Compute the rmse: rmse
rmse = np.sqrt(mean_squared_error(y_test, preds))
print("RMSE: %f" % (rmse))
```

RMSE: 78847.401758

```
In [19]: # Linear base learners

# Convert the training and testing sets into DMatrixes: DM_train, DM_test
DM_train = xgb.DMatrix(data=X_train,label=y_train)
DM_test = xgb.DMatrix(data=X_test,label=y_test)

# Create the parameter dictionary: params
params = {"booster":"gblinear", "objective":"reg:linear"}

# Train the model: xg_reg
xg_reg = xgb.train(dtrain= DM_train, params=params, num_boost_round=5)

# Predict the labels of the test set: preds
preds = xg_reg.predict(DM_test)

# Compute and print the RMSE
rmse = np.sqrt(mean_squared_error(y_test,preds))
print("RMSE: %f" % (rmse))

RMSE: 43566.535658
```

```
In [20]: # Evaluating model quality

# RMSE Model
# Create the DMatrix: housing_dmatrix
housing_dmatrix = xgb.DMatrix(data=X, label=y)

# Create the parameter dictionary: params
params = {"objective":"reg:linear", "max_depth":4}

# Perform cross-validation: cv_results
cv_results = xgb.cv(dtrain=housing_dmatrix, params=params, nfold=4, num_boost_round=5, metrics='rmse', as_pandas=True, seed=123)

# Print cv_results
print(cv_results)

# Extract and print final boosting round metric
print((cv_results["test-rmse-mean"]).tail(1))

      test-rmse-mean  test-rmse-std  train-rmse-mean  train-rmse-std
0    142980.433594    1193.791602    141767.531250     429.454591
1    104891.394532    1223.158855     102832.544922     322.469930
2     79478.937500    1601.344539     75872.615235     266.475960
3     62411.920899    2220.150028     57245.652344     273.625086
4     51348.279297    2963.377719     44401.298828     316.423666
4     51348.279297
Name: test-rmse-mean, dtype: float64
```

```
In [21]: # Create the DMatrix: housing_dmatrix
housing_dmatrix = xgb.DMatrix(data=X, label=y)

# Create the parameter dictionary: params
params = {"objective":"reg:linear", "max_depth":4}

# Perform cross-validation: cv_results
cv_results = xgb.cv(dtrain=housing_dmatrix, params=params, nfold=4, num_boost_round=5, metrics='mae', as_pandas=True, seed=123)

# Print cv_results
print(cv_results)

# Extract and print final boosting round metric
print((cv_results["test-mae-mean"]).tail(1))

      test-mae-mean  test-mae-std  train-mae-mean  train-mae-std
0    127634.000000    2404.009898    127343.482422     668.308109
1     90122.501954    2107.912810     89770.056641     456.965267
2     64278.558594    1887.567576     63580.791016     263.404950
3     46819.168946    1459.818607     45633.155274     151.883420
4     35670.646485    1140.607452     33587.090820      86.999396
4     35670.646485
Name: test-mae-mean, dtype: float64
```

Regularization and base learners in XGBoost

Regularization

- Regularization is a control on model complexity
- Want models that are both accurate and as simple as possible
- Tweak parameters to limit model complexity by altering loss function
- Regularization parameters in XGBoost:
 - gamma
 - for tree base learners
 - controls whether a node on base learner split based on the expected reduction in the loss that would occur after performing the split
 - higher values lead to fewer splits
 - minimum loss reduction allowed for a split to occur
 - alpha
 - another name for l1 regularization

- penalty on leaf weights rather than feature weights, as is the case in linear or logistic regression
- higher alpha values mean more regularization, causes many leaf weights in the base learners to be 0
- lambda
 - another name for l2 regularization
 - much smoother penalty than l1
 - causes leaf weights to smoothly decrease instead of enforcing strong sparsity constraints on the leaf weights as in l1.

```
In [22]: # L1 Regularization in XGBoost example

boston_data = pd.read_csv("datasets/boston.csv")
X,y = boston_data.iloc[:, :-1], boston_data.iloc[:, -1]

boston_dmatrix = xgb.DMatrix(data=X, label=y)
params={"objective": "reg:linear", "max_depth": 4}

l1_params = [1, 10, 100]
rmse_l1 = []

for reg in l1_params:
    params["alpha"] = reg
    cv_results = xgb.cv(dtrain=boston_dmatrix, params=params,
                        nfold=4, num_boost_round=10,
                        metrics="rmse", as_pandas=True, seed=123)
    rmse_l1.append(cv_results["test-rmse-mean"] \
                  .tail(1).values[0])
```

```
In [23]: print("Best rmse as a function of l1:")
print(pd.DataFrame(list(zip(l1_params, rmse_l1)), columns=["l1", "rmse"]))
```

Best rmse as a function of l1:

| | l1 | rmse |
|---|-----|----------|
| 0 | 1 | 3.461474 |
| 1 | 10 | 3.821152 |
| 2 | 100 | 4.645519 |

Base learners in XGBoost

- Linear Base Learner:
 - Sum of linear terms (same as in linear/logistic regression models)
 - When combined into an ensemble, the boosted model is weighted sum of linear models (thus is itself linear)
 - Don't get any nonlinear combination of features in the final model
 - Rarely used
- Tree Base Learner:
 - Decision trees as base model
 - When combined into an ensemble, boosted model is weighted sum of decision trees (nonlinear)
 - Almost exclusively used in XGBoost

```
In [24]: # EXERCISES
```

```
In [25]: # Using L2 regularization in XGBoost

# Create the DMatrix: housing_dmatrix
housing_dmatrix = xgb.DMatrix(data=X, label=y)

reg_params = [1, 10, 100]

# Create the initial parameter dictionary for varying L2 strength: params
params = {"objective": "reg:linear", "max_depth": 3}

# Create an empty list for storing rmse as a function of L2 complexity
rmse_l2 = []

# Iterate over reg_params
for reg in reg_params:

    # Update L2 strength
    params["lambda"] = reg

    # Pass this updated param dictionary into cv
    cv_results_rmse = xgb.cv(dtrain=housing_dmatrix, params=params, nfold=2, num_boost_round=5, metrics="rmse", as_pandas=True, seed=123)

    # Append best rmse (final round) to rmse_l2
    rmse_l2.append(cv_results_rmse["test-rmse-mean"].tail(1).values[0])

# Look at best rmse per L2 param
print("Best rmse as a function of l2:")
print(pd.DataFrame(list(zip(reg_params, rmse_l2)), columns=["l2", "rmse"]))
```

Best rmse as a function of l2:

| | l2 | rmse |
|---|-----|-----------|
| 0 | 1 | 6.022222 |
| 1 | 10 | 7.201520 |
| 2 | 100 | 10.692151 |


```
In [26]: # Visualizing individual XGBoost trees

# Create the DMatrix: housing_dmatrix
housing_dmatrix = xgb.DMatrix(data=X, label=y)

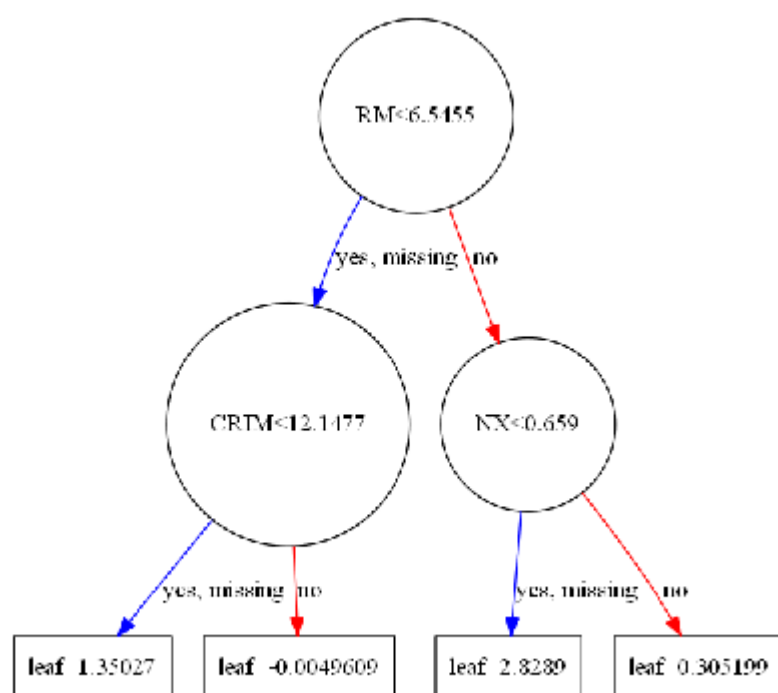
# Create the parameter dictionary: params
params = {"objective":"reg:linear", "max_depth":2}

# Train the model: xg_reg
xg_reg = xgb.train(params=params, dtrain=housing_dmatrix, num_boost_round=10)


# Plot the first tree
xgb.plot_tree(xg_reg, num_trees=0)
plt.rcParams["figure.figsize"] = [7,7]
plt.show()
```



```
In [27]: # Plot the fifth tree
xgb.plot_tree(xg_reg, num_trees=4)
plt.rcParams["figure.figsize"] = [7,7]
plt.show()
```



```
In [28]: # Plot the last tree sideways
xgb.plot_tree(xg_reg, num_trees=9, rankdir="LR")
plt.rcParams["figure.figsize"] = [10,10]
plt.show()
```





```
In [29]: # Visualizing feature importances: What features are most important in my dataset
```

```
# Create the DMatrix: housing_dmatrix
df = pd.read_csv('datasets/ames_housing_trimmed_processed.csv')
X, y = df.iloc[:, :-1], df.iloc[:, -1]
housing_dmatrix = xgb.DMatrix(data=X, label=y)

# Create the parameter dictionary: params
params = {'objective': 'reg:linear',
          'max_depth': 4
        }

# Train the model: xg_reg
xg_reg = xgb.train(dtrain=housing_dmatrix, params=params, num_boost_round=10)

# Plot the feature importances
xgb.plot_importance(xg_reg)
plt.show()
```



Chap 3: Fine-tuning your XGBoost model

```
In [30]: # Import plotting modules
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
plt.style.use('ggplot')

from sklearn.model_selection import train_test_split
import xgboost as xgb
from sklearn.metrics import mean_squared_error
```

Why tune your model?

```
In [31]: # Untuned Model Example

housing_data = pd.read_csv("datasets/ames_housing_trimmed_processed.csv")
X,y = housing_data[housing_data.columns.tolist()[:-1]],housing_data[housing_data.columns.tolist()[-1]]

housing_dmatrix = xgb.DMatrix(data=X,label=y)
untuned_params={"objective":"reg:linear"}

untuned_cv_results_rmse = xgb.cv(dtrain=housing_dmatrix,params=untuned_params,
                                nfold=4,metrics="rmse",as_pandas=True,seed=123)

print("Untuned rmse: %f" %((untuned_cv_results_rmse["test-rmse-mean"]).tail(1)))
```

Untuned rmse: 34624.229980

```
In [32]: # Tuned Model Example

housing_dmatrix = xgb.DMatrix(data=X,label=y)
tuned_params = {"objective":"reg:linear",'colsample_bytree': 0.3,
                'learning_rate': 0.1, 'max_depth': 5}

tuned_cv_results_rmse = xgb.cv(dtrain=housing_dmatrix,params=tuned_params,
                              nfold=4, num_boost_round=200, metrics="rmse",
                              as_pandas=True, seed=123)

print("Tuned rmse: %f" %((tuned_cv_results_rmse["test-rmse-mean"]).tail(1)))

Tuned rmse: 30187.115723
```

```
In [33]: # EXERCISES
```

```
In [34]: # Tuning the number of boosting rounds

# Create the DMatrix: housing_dmatrix
housing_dmatrix = xgb.DMatrix(data=X, label=y)

# Create the parameter dictionary for each tree: params
params = {"objective":"reg:linear", "max_depth":3}

# Create list of number of boosting rounds
num_rounds = [5, 10, 15]

# Empty list to store final round rmse per XGBoost model
final_rmse_per_round = []

# Iterate over num_rounds and build one model per num_boost_round parameter
for curr_num_rounds in num_rounds:

    # Perform cross-validation: cv_results
    cv_results = xgb.cv(dtrain=housing_dmatrix, params=params, nfold=3, num_boost_round=curr_num_rounds, metrics="rmse", as_pandas=

    # Append final round RMSE
    final_rmse_per_round.append(cv_results["test-rmse-mean"].tail(1).values[0])

# Print the resultant DataFrame
num_rounds_rmse = list(zip(num_rounds, final_rmse_per_round))
print(pd.DataFrame(num_rounds_rmse,columns=["num_boosting_rounds","rmse"]))
```

| | num_boosting_rounds | rmse |
|---|---------------------|--------------|
| 0 | 5 | 50903.299479 |
| 1 | 10 | 34774.194011 |
| 2 | 15 | 32895.098958 |

```
In [35]: # Automated boosting round selection using early_stopping

# Create your housing DMatrix: housing_dmatrix
housing_dmatrix = xgb.DMatrix(data=X, label=y)

# Create the parameter dictionary for each tree: params
params = {"objective":"reg:linear", "max_depth":4}

# Perform cross-validation with early stopping: cv_results
cv_results = xgb.cv(dtrain=housing_dmatrix,params=params,num_boost_round=50,nfold=3,metrics='rmse',early_stopping_rounds=10,as_pand

# Print cv_results
print(cv_results.tail())
```

| | test-rmse-mean | test-rmse-std | train-rmse-mean | train-rmse-std |
|----|----------------|---------------|-----------------|----------------|
| 45 | 30758.543620 | 1947.454953 | 11356.552734 | 565.368794 |
| 46 | 30729.971354 | 1985.698867 | 11193.557943 | 552.299272 |
| 47 | 30732.662760 | 1966.997355 | 11071.315755 | 604.090310 |
| 48 | 30712.241537 | 1957.751573 | 10950.778320 | 574.862779 |
| 49 | 30720.854167 | 1950.511057 | 10824.865560 | 576.665674 |

Overview of XGBoost's hyperparameters

Common TREE tunable parameters

- eta:
 - learning rate - value between 0 and 1
 - affects how quickly the model fits the residual error using additional base learners
 - low learning rate will require more boosting rounds to achieve the same reduction in residual error as an XGBoost model with high learning rate
- gamma:
 - min loss reduction to create new tree split
- lambda:
 - L2 reg on leaf weights
- alpha:
 - L1 reg on leaf weights
- max_depth:

- max depth per tree
- must be a positive integer value
- affects how deeply each tree is allowed to grow during any given boosting round
- subsample:
 - % samples used per tree
 - must be a value between 0 and 1
 - fraction of the total training set that can be used for any given boosting round
 - if value is low, fraction of training data used per boosting round would be low, causing underfitting problems.
 - if value is very high, can lead to overfitting
- colsample_bytree:
 - % features used per tree
 - fraction of features that you can select from during any given boosting round
 - same as RandomForest 'max_features' attribute
 - must also be a value between 0 and 1
 - large value means almost all features can be used to build a tree during a boosting
 - may in certain cases overfit a trained model
 - small value means that the fraction of features that can be selected from is very small.
 - smaller values can be thought of as providing additional regularization to the model
- n_estimators / no. of boosting rounds is tunable in both models
 - either no. of trees to build
 - or no. of linear base learners to construct

LINEAR tunable parameters

- no. of tunable parameters is significantly smaller
- lambda:
 - L2 reg on weights associated with any given feature
- alpha:
 - L1 reg on weights
- lambda_bias:
 - L2 reg term on model's bias

In [36]: `# EXERCISES`

In [37]: `# Tuning eta`

```
# Create your housing DMatrix: housing_dmatrix
housing_dmatrix = xgb.DMatrix(data=X, label=y)

# Create the parameter dictionary for each tree (boosting round)
params = {"objective":"reg:linear", "max_depth":3}

# Create list of eta values and empty list to store final round rmse per xgboost model
eta_vals = [0.001, 0.01, 0.1]
best_rmse = []

# Systematically vary the eta
for curr_val in eta_vals:
    params["eta"] = curr_val
    # Perform cross-validation: cv_results
    cv_results = xgb.cv(dtrain=housing_dmatrix,params=params,nfold=3,
                        num_boost_round=10,early_stopping_rounds=5,
                        metrics='rmse',seed=123)

    # Append the final round rmse to best_rmse
    best_rmse.append(cv_results["test-rmse-mean"].tail().values[-1])

# Print the resultant DataFrame
print(pd.DataFrame(list(zip(eta_vals, best_rmse)), columns=["eta","best_rmse"]))
```

| | eta | best_rmse |
|---|-------|---------------|
| 0 | 0.001 | 195736.406250 |
| 1 | 0.010 | 179932.182292 |
| 2 | 0.100 | 79759.411459 |

```
In [38]: # Tuning max_depth

# Create your housing DMatrix
housing_dmatrix = xgb.DMatrix(data=X,label=y)

# Create the parameter dictionary
params = {"objective":"reg:linear"}

# Create list of max_depth values
max_depths = [2,5,10,20]
best_rmse = []

# Systematically vary the max_depth
for curr_val in max_depths:
    params["max_depth"] = curr_val
    # Perform cross-validation
    cv_results = xgb.cv(dtrain=housing_dmatrix,params=params,metrics='rmse',
                        num_boost_round=10,early_stopping_rounds=5,
                        nfold=2,as_pandas=True,seed=123)

    # Append the final round rmse to best_rmse
    best_rmse.append(cv_results["test-rmse-mean"].tail().values[-1])

# Print the resultant DataFrame
print(pd.DataFrame(list(zip(max_depths, best_rmse)),columns=["max_depth","best_rmse"]))
```

| | max_depth | best_rmse |
|---|-----------|--------------|
| 0 | 2 | 37957.468750 |
| 1 | 5 | 35596.599610 |
| 2 | 10 | 36065.546875 |
| 3 | 20 | 36739.576172 |

```
In [39]: # Tuning colsample_bytree

# Create your housing DMatrix
housing_dmatrix = xgb.DMatrix(data=X,label=y)

# Create the parameter dictionary
params={"objective":"reg:linear","max_depth":3}

# Create list of hyperparameter values: colsample_bytree_vals
colsample_bytree_vals = [0.1,0.5,0.8,1]
best_rmse = []

# Systematically vary the hyperparameter value
for curr_val in colsample_bytree_vals:

    params['colsample_bytree'] = curr_val

    # Perform cross-validation
    cv_results = xgb.cv(dtrain=housing_dmatrix, params=params, nfold=2,
                        num_boost_round=10, early_stopping_rounds=5,
                        metrics="rmse", as_pandas=True, seed=123)

    # Append the final round rmse to best_rmse
    best_rmse.append(cv_results["test-rmse-mean"].tail().values[-1])

# Print the resultant DataFrame
print(pd.DataFrame(list(zip(colsample_bytree_vals, best_rmse)), columns=["colsample_bytree","best_rmse"]))
```

| | colsample_bytree | best_rmse |
|---|------------------|--------------|
| 0 | 0.1 | 51764.712890 |
| 1 | 0.5 | 35612.806641 |
| 2 | 0.8 | 35509.833985 |
| 3 | 1.0 | 35836.046875 |

Review of Grid Search and Random Search

- Hyperparameter values interact in non-obvious/non-linear ways
- Two common search strategies are:
 - Grid search
 - Random search
- both can be used with XGBoost and scikit-learn packages

Grid Search: Review

- Search exhaustively over a given set of hyperparameters, once per set of hyperparameters
- Number of models = number of distinct values per hyperparameter multiplied across each hyperparameter
 - for 2 hyperparameters to tune with 4 possible values for each, total number of models tried by grid search will be 16
- Pick final model hyperparameter values that give best cross-validated evaluation metric value

In [40]: *# Grid Search: Example*

```
from sklearn.model_selection import GridSearchCV

housing_data = pd.read_csv("datasets/ames_housing_trimmed_processed.csv")
X, y = housing_data[housing_data.columns.tolist()[:-1]], housing_data[housing_data.columns.tolist()[-1]]
housing_dmatrix = xgb.DMatrix(data=X,label=y)

gbm_param_grid = {'learning_rate': [0.01,0.1,0.5,0.9],
                  'n_estimators': [200],
                  'subsample': [0.3, 0.5, 0.9]}

gbm = xgb.XGBRegressor()
grid_mse = GridSearchCV(estimator=gbm, cv=4, verbose=1,
                       param_grid=gbm_param_grid,
                       scoring='neg_mean_squared_error')

grid_mse.fit(X, y)
```

Fitting 4 folds for each of 12 candidates, totalling 48 fits

[Parallel(n_jobs=1)]: Done 48 out of 48 | elapsed: 18.8s finished

Out[40]: GridSearchCV(cv=4, error_score='raise',
 estimator=XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
 colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,
 max_depth=3, min_child_weight=1, missing=None, n_estimators=100,
 n_jobs=1, nthread=None, objective='reg:linear', random_state=0,
 reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
 silent=True, subsample=1),
 fit_params=None, iid=True, n_jobs=1,
 param_grid={'learning_rate': [0.01, 0.1, 0.5, 0.9], 'n_estimators': [200], 'subsample': [0.3, 0.5, 0.9]},
 pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
 scoring='neg_mean_squared_error', verbose=1)

In [41]: print("Best parameters found: ", grid_mse.best_params_)
 print("Lowest RMSE found: ", np.sqrt(np.abs(grid_mse.best_score_)))

Best parameters found: {'learning_rate': 0.1, 'n_estimators': 200, 'subsample': 0.5}
 Lowest RMSE found: 28574.9861732

Random Search: Review

- Significantly different from Grid search
- No. of models required to iterate over doesn't grow as you expand the overall hyperparameter space
- Create a (possibly infinite) range of hyperparameter values per hyperparameter that you would like to search over
- Set the number of iterations you would like for the random search to continue
- During each iteration, randomly draw a value in the range of specified values for each hyperparameter searched over and train/evaluate a model with those hyperparameters
- After you've reached the maximum number of iterations, select the hyperparameter configuration with the best evaluated score

In [42]: *# Random Search: Example*

```
from sklearn.model_selection import RandomizedSearchCV

housing_data = pd.read_csv("datasets/ames_housing_trimmed_processed.csv")
X,y = housing_data[housing_data.columns.tolist()[:-1]],housing_data[housing_data.columns.tolist()[-1]]
housing_dmatrix = xgb.DMatrix(data=X,label=y)

gbm_param_grid = {'learning_rate': np.arange(0.05,1.05,.05),
                  'n_estimators': [200],
                  'subsample': np.arange(0.05,1.05,.05)}

gbm = xgb.XGBRegressor()
randomized_mse = RandomizedSearchCV(estimator=gbm, n_iter=25,
                                    scoring='neg_mean_squared_error',
                                    param_distributions=gbm_param_grid,
                                    cv=4, verbose=1)

randomized_mse.fit(X, y)
```

Fitting 4 folds for each of 25 candidates, totalling 100 fits

[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 41.1s finished

Out[42]: RandomizedSearchCV(cv=4, error_score='raise',
 estimator=XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
 colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,
 max_depth=3, min_child_weight=1, missing=None, n_estimators=100,
 n_jobs=1, nthread=None, objective='reg:linear', random_state=0,
 reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
 silent=True, subsample=1),
 fit_params=None, iid=True, n_iter=25, n_jobs=1,
 param_distributions={'learning_rate': array([0.05, 0.1 , 0.15, 0.2 , 0.25, 0.3 , 0.35, 0.4 , 0.45,
 0.5 , 0.55, 0.6 , 0.65, 0.7 , 0.75, 0.8 , 0.85, 0.9 ,
 0.95, 1.]), 'n_estimators': [200], 'subsample': array([0.05, 0.1 , 0.15, 0.2 , 0.25, 0.3 , 0.35, 0.4 , 0.45,
 0.5 , 0.55, 0.6 , 0.65, 0.7 , 0.75, 0.8 , 0.85, 0.9 ,
 0.95, 1.])},
 pre_dispatch='2*n_jobs', random_state=None, refit=True,
 return_train_score='warn', scoring='neg_mean_squared_error',
 verbose=1)

```
In [43]: print("Best parameters found: ", randomized_mse.best_params_)
print("Lowest RMSE found: ", np.sqrt(np.abs(randomized_mse.best_score_)))

Best parameters found: {'subsample': 0.35000000000000003, 'n_estimators': 200, 'learning_rate': 0.15000000000000002}
Lowest RMSE found: 28753.2590098

In [44]: # EXERCISES

In [45]: # Grid Search with XGBoost

# Create your housing DMatrix: housing_dmatrix
housing_dmatrix = xgb.DMatrix(data=X, label=y)

# Create the parameter grid: gbm_param_grid
gbm_param_grid = {
    'colsample_bytree': [0.3, 0.7],
    'n_estimators': [50],
    'max_depth': [2, 5]
}

# Instantiate the regressor: gbm
gbm = xgb.XGBRegressor()

# Perform grid search: grid_mse
grid_mse = GridSearchCV(estimator=gbm, param_grid=gbm_param_grid, cv=4, scoring='neg_mean_squared_error', verbose=1)

# Fit grid_mse to the data
grid_mse.fit(X, y)

# Print the best parameters and Lowest RMSE
print("Best parameters found: ", grid_mse.best_params_)
print("Lowest RMSE found: ", np.sqrt(np.abs(grid_mse.best_score_)))
```

Fitting 4 folds for each of 4 candidates, totalling 16 fits
Best parameters found: {'colsample_bytree': 0.3, 'max_depth': 5, 'n_estimators': 50}
Lowest RMSE found: 30031.6171206

[Parallel(n_jobs=1)]: Done 16 out of 16 | elapsed: 1.2s finished

```
In [46]: # Random Search with XGBoost

# Create the parameter grid: gbm_param_grid
gbm_param_grid = {
    'n_estimators': [25],
    'max_depth': range(2, 12)
}

# Instantiate the regressor: gbm
gbm = xgb.XGBRegressor(n_estimators=10)

# Perform random search: grid_mse
randomized_mse = RandomizedSearchCV(estimator=gbm, cv=4, n_iter=5, verbose=1,
                                    param_distributions=gbm_param_grid,
                                    scoring='neg_mean_squared_error')

# Fit randomized_mse to the data
randomized_mse.fit(X, y)

# Print the best parameters and Lowest RMSE
print("Best parameters found: ", randomized_mse.best_params_)
print("Lowest RMSE found: ", np.sqrt(np.abs(randomized_mse.best_score_)))
```

Fitting 4 folds for each of 5 candidates, totalling 20 fits

[Parallel(n_jobs=1)]: Done 20 out of 20 | elapsed: 2.4s finished

Best parameters found: {'n_estimators': 25, 'max_depth': 11}
Lowest RMSE found: 37502.1924786

Limits of Grid Search and Random Search

Grid Search

- Number of models you must build with every additional new parameter grows very quickly
- Optimal if no. of hyperparameters and distinct values per hyperparameter is kept small
- Time taken increases exponentially

Random Search

- Parameter space to explore can be massive
- Randomly jumping throughout the space looking for a "best" result becomes a waiting game

Chap 4: Using XGBoost in pipelines

```
In [87]: # Import plotting modules
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
plt.style.use('ggplot')
```

Review of pipelines using sklearn

Pipeline Review

- Takes a list of named 2-tuples (name, pipeline_step) as input
- Tuples can contain any arbitrary scikit-learn compatible estimator or transformer object
- Pipeline implements fit/predict methods
- Can be used as input estimator into methods like
 - grid/randomized search approaches - for tuning hyperparameters
 - cross_val_score - for efficient cross-validation and out of sample metric calculation

```
In [67]: from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.model_selection import cross_val_score

names = ["crime","zone","industry","charles","no","rooms","age",
         "distance","radial","tax","pupil","aam","lower","med_price"]

data = pd.read_csv("datasets/boston.csv",names=names,skiprows=1)
X, y = data.iloc[:, :-1], data.iloc[:, -1]

rf_pipeline = Pipeline([("st_scaler",StandardScaler()),
                        ("rf_model",RandomForestRegressor())])

scores = cross_val_score(rf_pipeline, X, y, cv=10,
                        scoring="neg_mean_squared_error")
```

```
In [68]: final_avg_rmse = np.mean(np.sqrt(np.abs(scores)))
print("Final RMSE:", final_avg_rmse)
```

Final RMSE: 4.30351262932

- neg_mean_squared_error is scikit-learn's way of calculating the mean squared error in an API-compatible way.
- Negative mean squared errors don't exist as all squares must be positive when working with real numbers.
- Thus we simply take the absolute value of the scores, take each of their square roots, and compute their mean to get root mean squared error across all 10 cross-validation folds.

Preprocessing

- Do the same preprocessing in two different ways
- only one of them can be done within a pipeline

First Approach - LabelEncoder and OneHotEncoder

- LabelEncoder: Converts a categorical column of strings into integers
- OneHotEncoder: Takes the column of integers and encodes them as dummy variables
- Cannot be done within a pipeline

Second Approach - DictVectorizer

- can accomplish both steps in one line of code
- found in feature extraction sub-module
- Traditionally used in text processing pipelines
- Converts lists of feature mappings into vectors
- Does not directly work with dataframes
 - Using pandas dataframes we don't initially have these features in lists form
 - Need to convert DataFrame into a list of dictionary entries

```
In [90]: # EXERCISES
```

```
In [71]: # Exploratory data analysis
# (of unprocessed Ames housing dataset)

df = pd.read_csv('datasets/ames_unprocessed_data.csv')
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 21 columns):
MSSubClass      1460 non-null int64
MSZoning        1460 non-null object
LotFrontage     1201 non-null float64
LotArea         1460 non-null int64
Neighborhood    1460 non-null object
BldgType        1460 non-null object
HouseStyle      1460 non-null object
OverallQual     1460 non-null int64
OverallCond     1460 non-null int64
YearBuilt       1460 non-null int64
Remodeled       1460 non-null int64
GrLivArea       1460 non-null int64
BsmtFullBath    1460 non-null int64
BsmtHalfBath    1460 non-null int64
FullBath        1460 non-null int64
HalfBath        1460 non-null int64
BedroomAbvGr    1460 non-null int64
Fireplaces      1460 non-null int64
GarageArea      1460 non-null int64
PavedDrive      1460 non-null object
SalePrice       1460 non-null int64
dtypes: float64(1), int64(15), object(5)
memory usage: 239.6+ KB
```

```
In [72]: # Encoding categorical columns I: LabelEncoder

# Import LabelEncoder
from sklearn.preprocessing import LabelEncoder

# Fill missing values with 0
df.LotFrontage = df.LotFrontage.fillna(0)

# Create a boolean mask for categorical columns
categorical_mask = (df.dtypes == 'object')

# Get list of categorical column names
categorical_columns = df.columns[categorical_mask].tolist()

# Print the head of the categorical columns
print(df[categorical_columns].head())

# Create LabelEncoder object: le
le = LabelEncoder()

# Apply LabelEncoder to categorical columns
df[categorical_columns] = df[categorical_columns].apply(lambda x: le.fit_transform(x))

# Print the head of the LabelEncoded categorical columns
print(df[categorical_columns].head())
```

| | MSZoning | Neighborhood | BldgType | HouseStyle | PavedDrive |
|---|----------|--------------|----------|------------|------------|
| 0 | RL | CollgCr | 1Fam | 2Story | Y |
| 1 | RL | Veenker | 1Fam | 1Story | Y |
| 2 | RL | CollgCr | 1Fam | 2Story | Y |
| 3 | RL | Crawfor | 1Fam | 2Story | Y |
| 4 | RL | NoRidge | 1Fam | 2Story | Y |
| | MSZoning | Neighborhood | BldgType | HouseStyle | PavedDrive |
| 0 | 3 | 5 | 0 | 5 | 2 |
| 1 | 3 | 24 | 0 | 2 | 2 |
| 2 | 3 | 5 | 0 | 5 | 2 |
| 3 | 3 | 6 | 0 | 5 | 2 |
| 4 | 3 | 15 | 0 | 5 | 2 |

In [73]:

```
# Import OneHotEncoder
from sklearn.preprocessing import OneHotEncoder

# Create OneHotEncoder: ohe
ohe = OneHotEncoder(categorical_features=categorical_mask,sparse=False)

# Apply OneHotEncoder to categorical columns - output is no longer a dataframe: df_encoded
df_encoded = ohe.fit_transform(df)

# Print first 5 rows of the resulting dataset - again, this will no longer be a pandas dataframe
print(df_encoded[:5, :])

# Print the shape of the original DataFrame
print(df.shape)

# Print the shape of the transformed array
print(df_encoded.shape)
```

[illegible]

```
0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00
0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00
1.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00
0.00000000e+00 1.00000000e+00 6.00000000e+01 8.40000000e+01
1.42600000e+04 8.00000000e+00 5.00000000e+00 2.00000000e+03
0.00000000e+00 2.19800000e+03 1.00000000e+00 0.00000000e+00
2.00000000e+00 1.00000000e+00 4.00000000e+00 1.00000000e+00
8.36000000e+02 2.50000000e+05]]
(1460, 21)
(1460, 62)
```

In [75]: `# Encoding categorical columns III: DictVectorizer`

```
# Import DictVectorizer
from sklearn.feature_extraction import DictVectorizer

# Convert df into a dictionary: df_dict
df_dict = df.to_dict('records')

# Create the DictVectorizer object: dv
dv = DictVectorizer()

# Apply dv on df: df_encoded
df_encoded = dv.fit_transform(df_dict)

# Print the resulting first five rows
print(df_encoded[:5,:])

# Print the vocabulary
print(dv.vocabulary_)
```

```
(0, 0)      3.0
(0, 2)      1.0
(0, 5)      2.0
(0, 6)     548.0
(0, 7)    1710.0
(0, 8)      1.0
(0, 9)      5.0
(0, 10)    8450.0
(0, 11)     65.0
(0, 12)     60.0
(0, 13)      3.0
(0, 14)      5.0
(0, 15)      5.0
(0, 16)      7.0
(0, 17)      2.0
(0, 19)   208500.0
(0, 20)    2003.0
(1, 0)      3.0
(1, 3)      1.0
(1, 4)      1.0
(1, 5)      2.0
(1, 6)     460.0
(1, 7)    1262.0
(1, 9)      2.0
(1, 10)    9600.0
:          :
(3, 14)      6.0
(3, 15)      5.0
(3, 16)      7.0
(3, 17)      2.0
(3, 18)      1.0
(3, 19)   140000.0
(3, 20)    1915.0
(4, 0)       4.0
(4, 2)       1.0
(4, 4)       1.0
(4, 5)       2.0
(4, 6)     836.0
(4, 7)    2198.0
(4, 8)       1.0
(4, 9)       5.0
(4, 10)   14260.0
(4, 11)     84.0
(4, 12)     60.0
(4, 13)      3.0
(4, 14)     15.0
(4, 15)      5.0
(4, 16)      8.0
(4, 17)      2.0
(4, 19)   250000.0
(4, 20)    2000.0
{'MSSubClass': 12, 'MSZoning': 13, 'LotFrontage': 11, 'LotArea': 10, 'Neighborhood': 14, 'BldgType': 1, 'HouseStyle': 9, 'OverallQual': 16, 'OverallCond': 15, 'YearBuilt': 20, 'Remodeled': 18, 'GrLivArea': 7, 'BsmtFullBath': 2, 'BsmtHalfBath': 3, 'FullBath': 5, 'HalfBath': 8, 'BedroomAbvGr': 0, 'Fireplaces': 4, 'GarageArea': 6, 'PavedDrive': 17, 'SalePrice': 19}
```

```
In [77]: # Preprocessing within a pipeline

# Import necessary modules
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction import DictVectorizer

X, y = df.iloc[:, :-1], df.iloc[:, -1]

# Fill LotFrontage missing values with 0
X.LotFrontage = X.LotFrontage.fillna(0)

# Setup the pipeline steps: steps
steps = [("ohe_onestep", DictVectorizer(sparse=False)),
        ("xgb_model", xgb.XGBRegressor())]

# Create the pipeline: xgb_pipeline
xgb_pipeline = Pipeline(steps)

# Fit the pipeline
xgb_pipeline.fit(X.to_dict('records'), y)
```

```
Out[77]: Pipeline(memory=None,
      steps=[('ohe_onestep', DictVectorizer(dtype=<class 'numpy.float64'>, separator='=', sort=True,
      sparse=False)), ('xgb_model', XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
      colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,
      max_depth=3, min_ch...
      reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
      silent=True, subsample=1))])
```

Incorporating XGBoost into pipelines

To have XGBoost in pipeline, use its scikit-learn API within a pipeline object.

```
In [80]: # Scikit-Learn Pipeline Example With XGBoost

from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.model_selection import cross_val_score

names = ["crime", "zone", "industry", "charles", "no", "rooms", "age",
        "distance", "radial", "tax", "pupil", "aam", "lower", "med_price"]
data = pd.read_csv("datasets/boston.csv", names=names, skiprows=1)
X, y = data.iloc[:, :-1], data.iloc[:, -1]

xgb_pipeline = Pipeline([("st_scaler", StandardScaler()),
                        ("xgb_model", xgb.XGBRegressor())])

scores = cross_val_score(xgb_pipeline, X, y, cv=10,
                        scoring="neg_mean_squared_error")
final_avg_rmse = np.mean(np.sqrt(np.abs(scores)))
print("Final XGB RMSE:", final_avg_rmse)
```

Final XGB RMSE: 4.02719593323

Additional Components Introduced For Pipelines

- sklearn_pandas: (sklearn and pandas both not always work seamlessly)
 - DataFrameMapper - Interoperability between pandas and scikit-learn
 - CategoricalImputer - Allow for imputation of categorical variables before conversion to integers
- sklearn.preprocessing:
 - Imputer - Native imputation of numerical columns in scikit-learn
- sklearn.pipeline:
 - FeatureUnion - combine multiple pipelines of features into a single pipeline of features

```
In [100]: # EXERCISES
```

```
In [82]: # Cross-validating your XGBoost model

# Import necessary modules
from sklearn.feature_extraction import DictVectorizer
from sklearn.pipeline import Pipeline
from sklearn.model_selection import cross_val_score

X, y = df.iloc[:, :-1], df.iloc[:, -1]

# Fill LotFrontage missing values with 0
X.LotFrontage = X.LotFrontage.fillna(0)

# Setup the pipeline steps: steps
steps = [("ohe_onestep", DictVectorizer(sparse=False)),
        ("xgb_model", xgb.XGBRegressor(max_depth=2, objective="reg:linear"))]

# Create the pipeline: xgb_pipeline
xgb_pipeline = Pipeline(steps)

# Cross-validate the model
cross_val_scores = cross_val_score(xgb_pipeline, X.to_dict('records'), y, cv=10, scoring='neg_mean_squared_error')

# Print the 10-fold RMSE
print("10-fold RMSE: ", np.mean(np.sqrt(np.abs(cross_val_scores))))

10-fold RMSE:  30343.4865518
```

```
In [100]: # Kidney disease case study I: Categorical Imputer

kidney_data = pd.read_csv('datasets/chronic_kidney_disease.csv', header=None, na_values='?')

kidney_feature_names = ['age', 'bp', 'sg', 'al', 'su', 'bgr', 'bu', 'sc', 'sod',
                        'pot', 'hemo', 'pcv', 'wc', 'rc', 'rbc', 'pc', 'pcc',
                        'ba', 'htn', 'dm', 'cad', 'appet', 'pe', 'ane']
kidney_target_name = ['class']
df.columns = kidney_feature_names + kidney_target_name
X, y = kidney_data.iloc[:, :-1], kidney_data.iloc[:, -1]
```

```
In [ ]: # Import necessary modules
from sklearn_pandas import DataFrameMapper
from sklearn_pandas import CategoricalImputer

# Check number of nulls in each feature column
nulls_per_column = X.isnull().sum()
print(nulls_per_column)

# Create a boolean mask for categorical columns
categorical_feature_mask = X.dtypes == object

# Get list of categorical column names
categorical_columns = X.columns[categorical_feature_mask].tolist()

# Get list of non-categorical column names
non_categorical_columns = X.columns[~categorical_feature_mask].tolist()

# Apply numeric imputer
numeric_imputation_mapper = \
DataFrameMapper([(numeric_feature, Imputer(strategy="median")) \
                 for numeric_feature in non_categorical_columns], \
                 input_df=True,
                 df_out=True)

# Apply categorical imputer
categorical_imputation_mapper = \
DataFrameMapper([(category_feature, CategoricalImputer()) \
                 for category_feature in categorical_columns], \
                 input_df=True,
                 df_out=True)
```

```
In [ ]: # Kidney disease case study II: Feature Union

# Import FeatureUnion
from sklearn.pipeline import FeatureUnion

# Combine the numeric and categorical transformations
numeric_categorical_union = \
FeatureUnion([("num_mapper", numeric_imputation_mapper), \
              ("cat_mapper", categorical_imputation_mapper)])
```

```
In [ ]: # Kidney disease case study III: Full pipeline
```

```
# Create full pipeline
pipeline = Pipeline([
    ("featureunion", numeric_categorical_union),
    ("dictifier", Dictifier()),
    ("vectorizer", DictVectorizer(sort=False)),
    ("clf", xgb.XGBClassifier(max_depth=3))
])

# Perform cross-validation
cross_val_scores = cross_val_score(pipeline, X, y, scoring="roc_auc", cv=3)

# Print avg. AUC
print("3-fold AUC: ", np.mean(cross_val_scores))
```

Tuning XGBoost hyperparameters

NOTE: Remember, in order to pass hyperparameters to the appropriate step,

- you have to name the parameters in the dictionary with the name of the step being referenced
- followed by 2 underscore signs and then the name of the hyperparameter you want to iterate over.

Since the xgboost step is called `xgb_model`, all of our hyperparameter keys will start with `xgboost_model__`.

```
In [105]: # Tuning XGBoost hyperparameters in a Pipeline
```

```
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.model_selection import RandomizedSearchCV

names = ["crime","zone","industry","charles","no","rooms","age",
         "distance","radial","tax","pupil","aam","lower","med_price"]
data = pd.read_csv("datasets/boston.csv",names=names,skiprows=1)
X, y = data.iloc[:, :-1], data.iloc[:, -1]

xgb_pipeline = Pipeline([("st_scaler", StandardScaler()),
                          ("xgb_model", xgb.XGBRegressor())])

gbm_param_grid = {'xgb_model__subsample': np.arange(.05, 1, .05),
                  'xgb_model__max_depth': np.arange(3, 20, 1),
                  'xgb_model__colsample_bytree': np.arange(.1, 1.05, .05) }

randomized_neg_mse = \
RandomizedSearchCV (estimator=xgb_pipeline, n_iter=10,
                    param_distributions=gbm_param_grid,
                    scoring='neg_mean_squared_error', cv=4)

randomized_neg_mse.fit(X, y)
```

```
Out[105]: RandomizedSearchCV(cv=4, error_score='raise',
                             estimator=Pipeline(memory=None,
                                                  steps=[('st_scaler', StandardScaler(copy=True, with_mean=True, with_std=True)), ('xgb_model', XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0,
max_depth=3, min_child_weight=1, missing=None,
reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
silent=True, subsample=1))]),
                             fit_params=None, iid=True, n_iter=10, n_jobs=1,
                             param_distributions={'xgb_model__subsample': array([ 0.05,  0.1 ,  0.15,  0.2 ,  0.25,  0.3 ,  0.35,  0.4 ,  0.45,
    0.5 ,  0.55,  0.6 ,  0.65,  0.7 ,  0.75,  0.8 ,  0.85,  0.9 ,  0.95]), 'xgb_model__max_depth': array([ 3,  4,  5,  6,  7,
    8,  9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19]), 'xgb_model__colsample_bytree': array([ 0.1 ,  0.15,  0.2 ,  0.25,  0.3 ,  0.35,
    0.4 ,  0.45,  0.5 ,
    0.55,  0.6 ,  0.65,  0.7 ,  0.75,  0.8 ,  0.85,  0.9 ,  0.95,  1. ])}},
                             pre_dispatch='2*n_jobs', random_state=None, refit=True,
                             return_train_score='warn', scoring='neg_mean_squared_error',
                             verbose=0)
```

```
In [107]: # Tuning XGBoost hyperparameters in a Pipeline II
```

```
print("Best rmse: ", np.sqrt(np.abs(randomized_neg_mse.best_score_)))
```

Best rmse: 4.47691694445

```
In [108]: print("Best model: ", randomized_neg_mse.best_estimator_)
```

```
Best model: Pipeline(memory=None,
                     steps=[('st_scaler', StandardScaler(copy=True, with_mean=True, with_std=True)), ('xgb_model', XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bytree=1.0000000000000004, gamma=0, learning_rate=0.1,
max_delta_step=0, max_depth=11, min_child_weight=1, missing=None,
reg_lambda=1, scale_pos_weight=1, seed=None, silent=True, subsample=0.7500000000000011))])
```

```
In [103]: # EXERCISES
```

```
In [ ]: # Bringing it all together

# Create the parameter grid
gbm_param_grid = {
    'clf__learning_rate': np.arange(0.05, 1, 0.05),
    'clf__n_estimators': np.arange(50, 200, 50),
    'clf__max_depth': np.arange(3, 10, 1)}

# Perform RandomizedSearchCV
randomized_roc_auc = RandomizedSearchCV(estimator=pipeline,param_distributions=gbm_param_grid,scoring='roc_auc',cv=2,n_iter=2,verbo

# Fit the estimator
randomized_roc_auc.fit(X,y)

# Compute metrics
print(randomized_roc_auc.best_score_)
print(randomized_roc_auc.best_estimator_)
```

Final Thoughts

What We Have Covered And You Have Learned

- Using XGBoost for classification tasks
- Using XGBoost for regression tasks
- Tuning XGBoost's most important hyperparameters
- Incorporating XGBoost into sklearn pipelines
- Advanced functions that allow us to seamlessly work with Pandas DataFrames and scikit-learn.

What We Have NOT COVERED (And How You Can Proceed)

- Using XGBoost for ranking/recommendation problems (Netflix/Amazon problem)
 - can be done by modifying the loss function we use when constructing the model
- Using more sophisticated hyperparameter tuning strategies for tuning XGBoost models (Bayesian Optimization)
 - entire companies have been created just for using this method in tuning models.
 - example: the company 'sigpot'
- Using XGBoost as part of an ensemble of other models for regression/classification
 - predictions we get from XGBoost model can be combined with other models
 - very powerful additional way to squeeze the last bit of juice from the data.

```
In [ ]:
```

```
In [ ]:
```